

# MLOps on Google Cloud

Training Course

## What this course is **NOT** about

- Data Science
- Designing / Building Models
- Math

## What this course is **about**

- Model Deployment / Serving
- Continuous (re)Training : CI/CD/CT
- Automation

# Machine Learning Overview

- Computer Vision (Images, Video)
  - Classification / Localization
  - Object Detection / Tracking
  - Segmentation

**Semantic Segmentation**



GRASS, CAT,  
TREE, SKY

No objects, just pixels

**Classification  
+ Localization**



CAT

Single Object

**Object  
Detection**



DOG, DOG, CAT

Multiple Object

**Instance  
Segmentation**

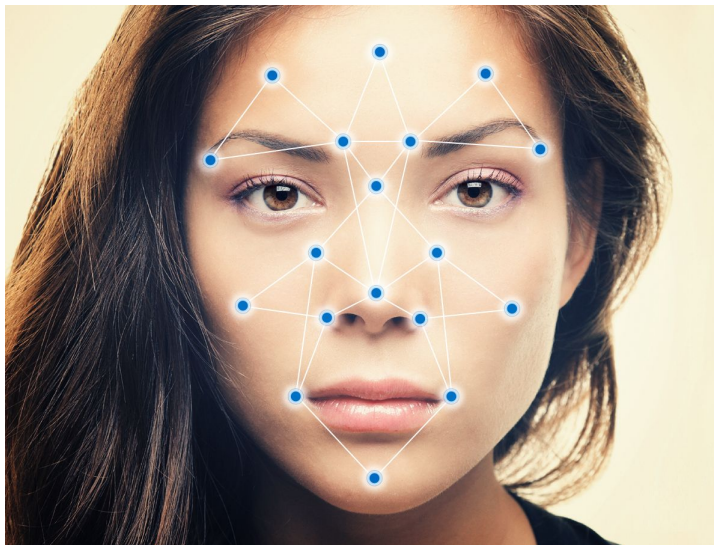


DOG, DOG, CAT

This image is CC0 public domain

# Machine Learning Overview

- Computer Vision (Images, Video)
  - Facial Recognition
  - Pose Detection
  - Captioning



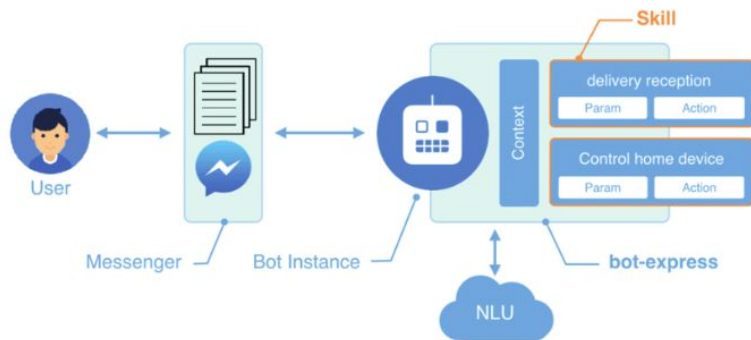
# Machine Learning Overview

- Natural Language Understanding (Text)
  - Classification
  - Sentiment
  - Entity Extraction
  - Form Recognition

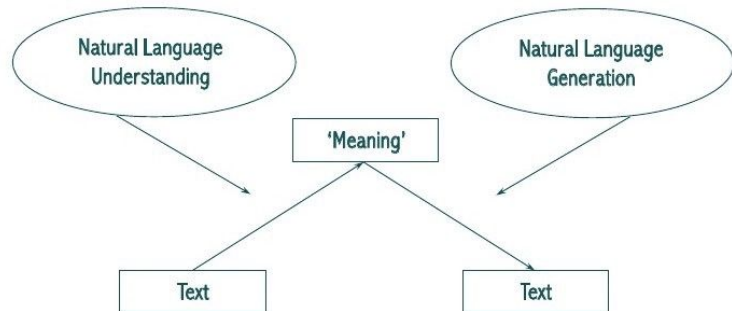


# Machine Learning Overview

- Natural Language Generation (Text/Audio)
  - Text-2-Speech / Speech-2-Text
  - Summarization
  - Chat

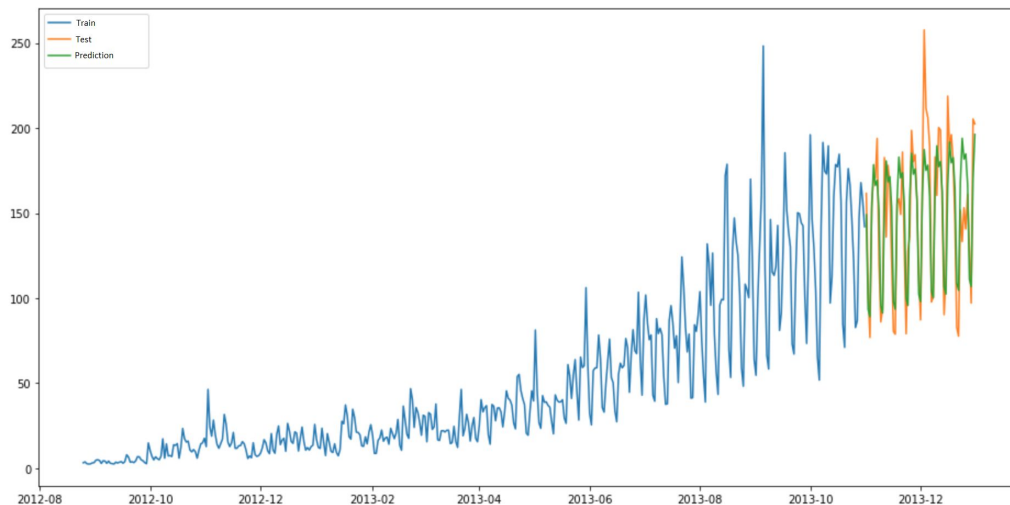
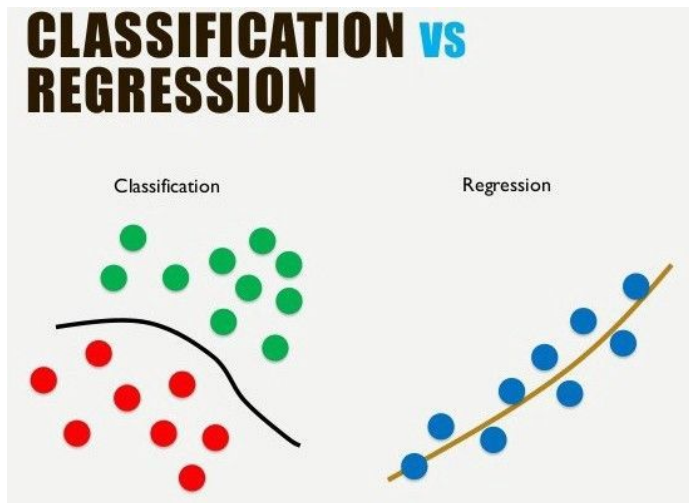


**NLP = NLU + NLG**



# Machine Learning Overview

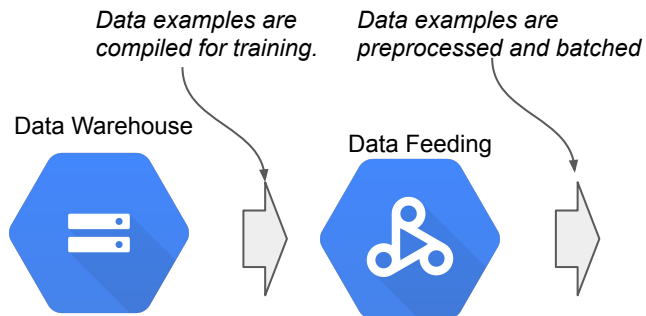
- Structured Data (Tabular, Databases)
  - Classification
  - Regression (Real Number)
  - Forecasting (time-series)





# MLOps Overview

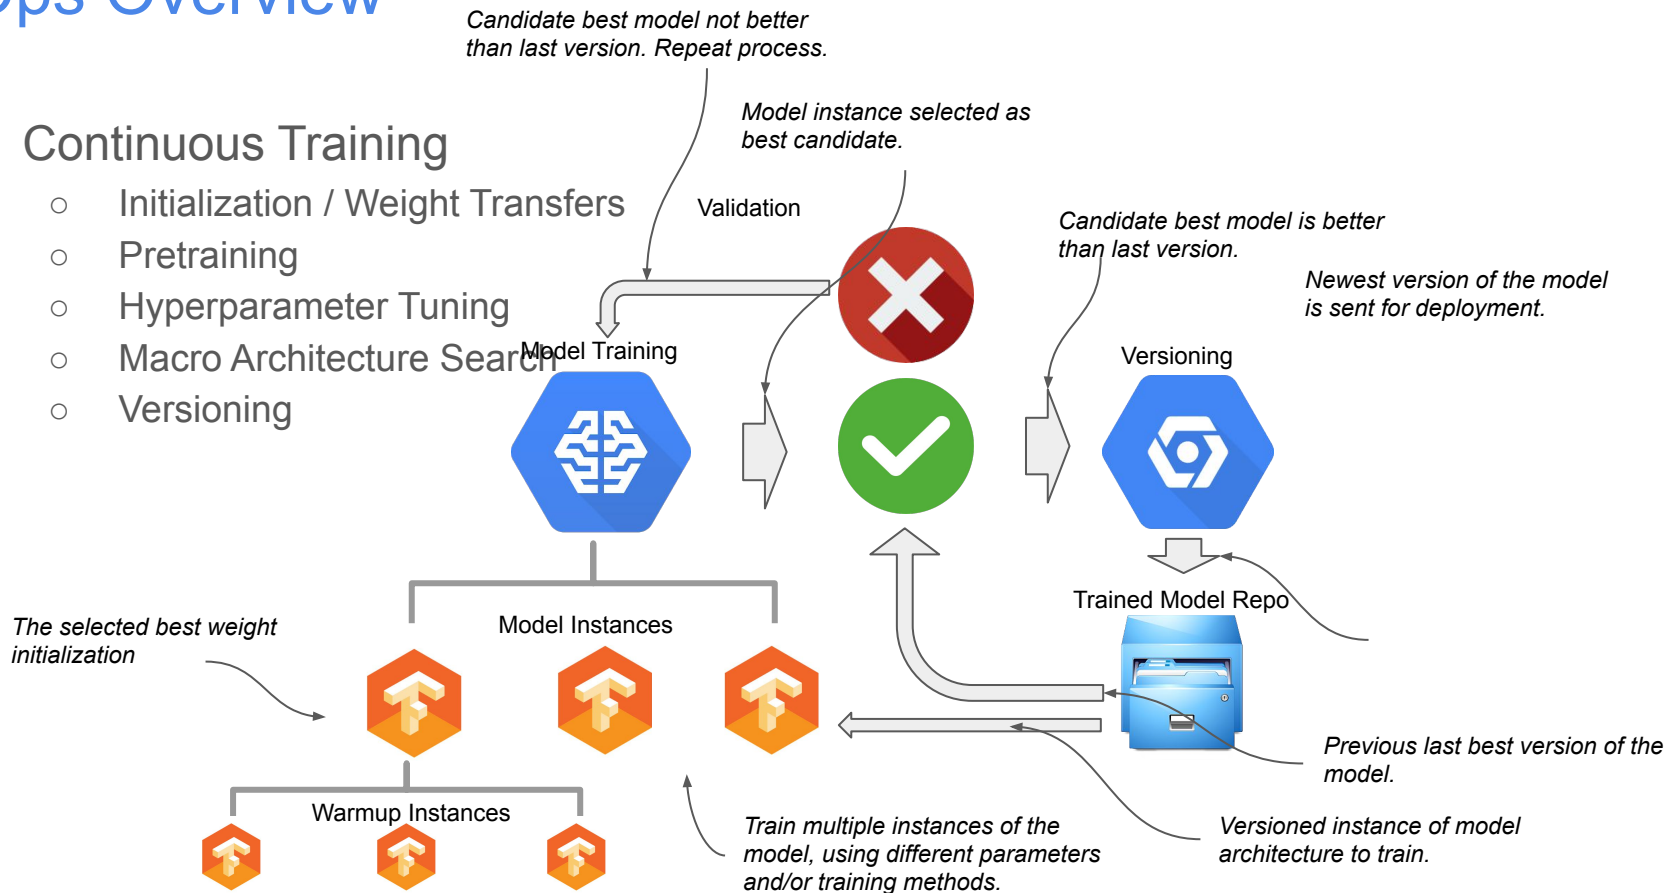
- Data Warehousing
  - Storage
  - Retrieval (I/O)
  - Feeding
  - Search / Query



# MLOps Overview

- Continuous Training

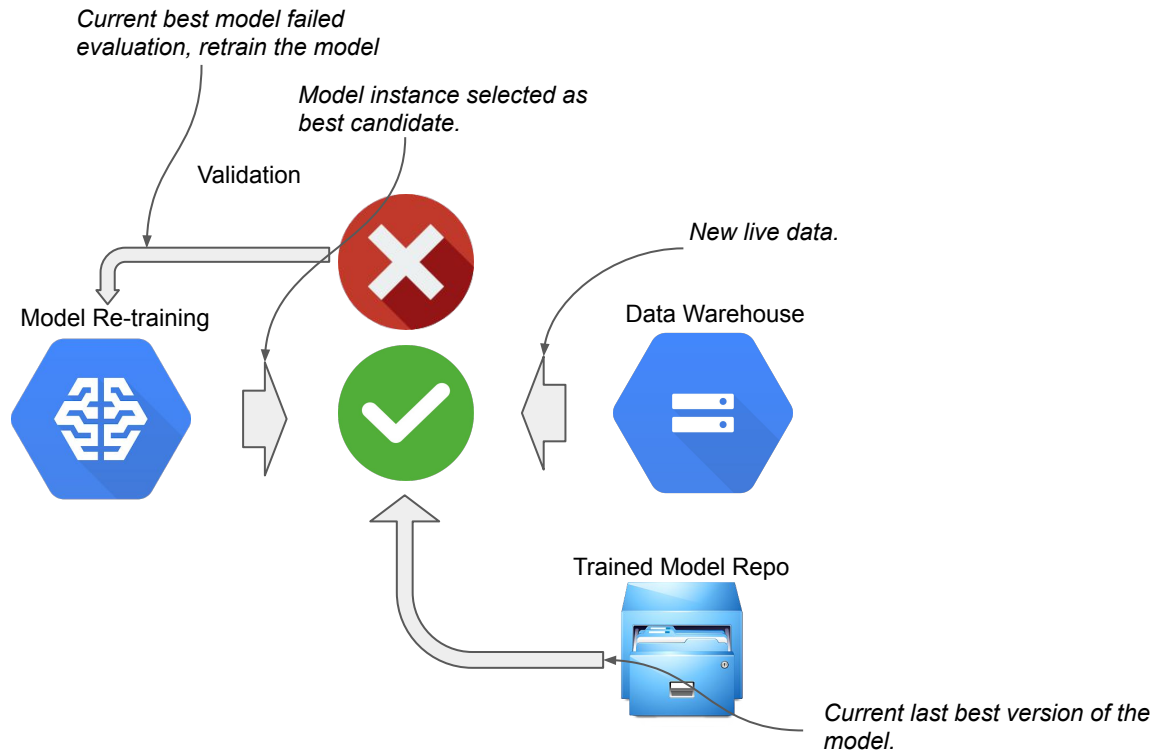
- Initialization / Weight Transfers
- Pretraining
- Hyperparameter Tuning
- Macro Architecture Search
- Versioning



# MLOps Overview

- Continuous Evaluation

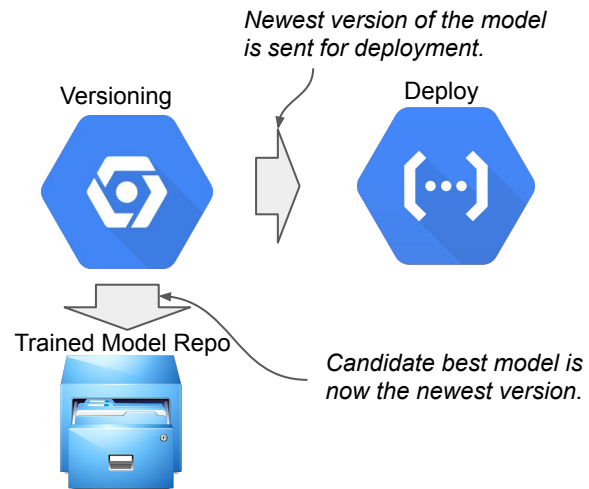
- Training Distribution
- Serving Skew
- Data Drift
- A/B Testing



# MLOps Overview

- Deployment

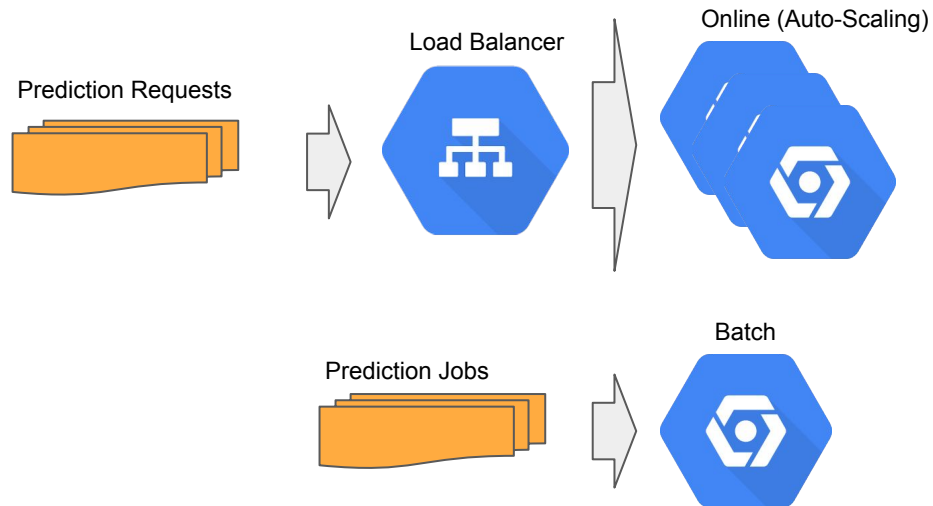
- Scaling
- Load Balancing
- Latency
- Edge



# MLOps Overview

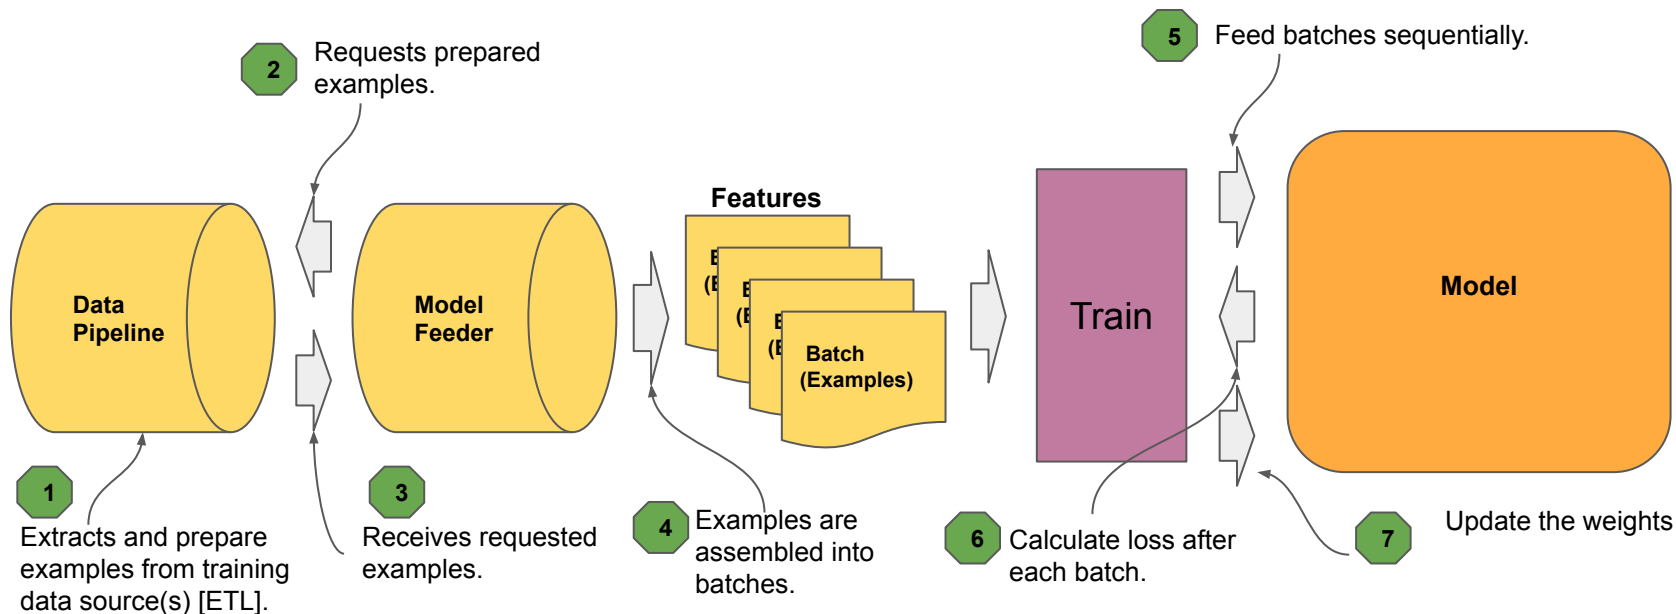
- Serving

- Online (live)
- Batch
- Monitoring
- Data Collection



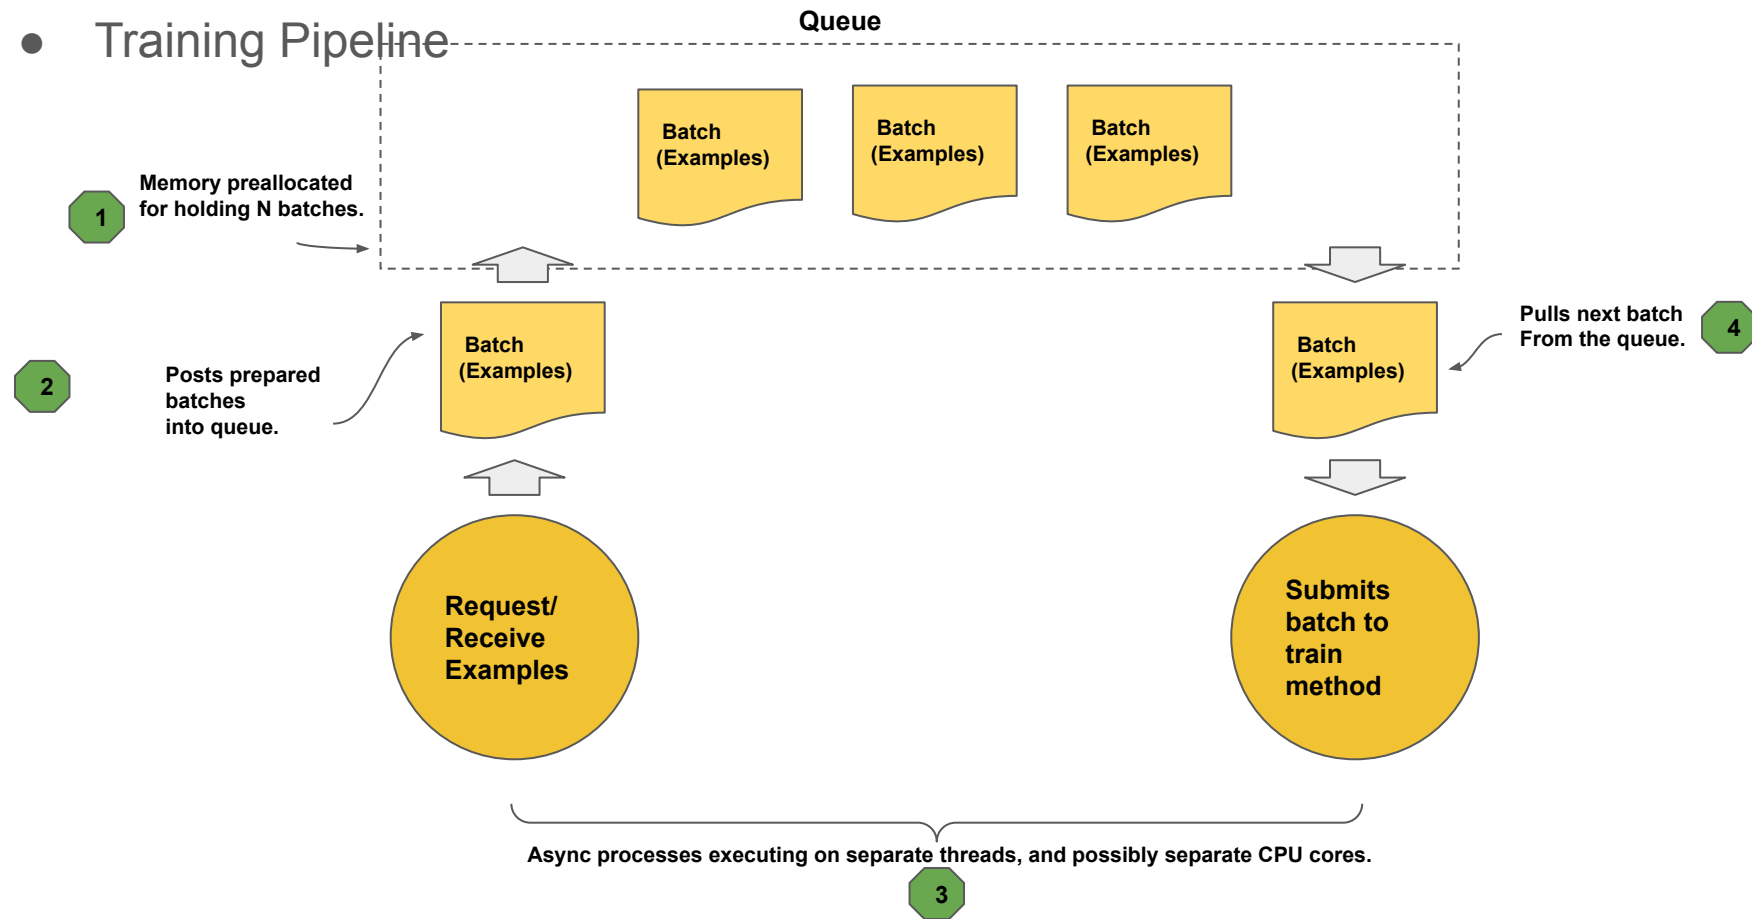
# ML end-2-end production pipeline

- Data Pipeline



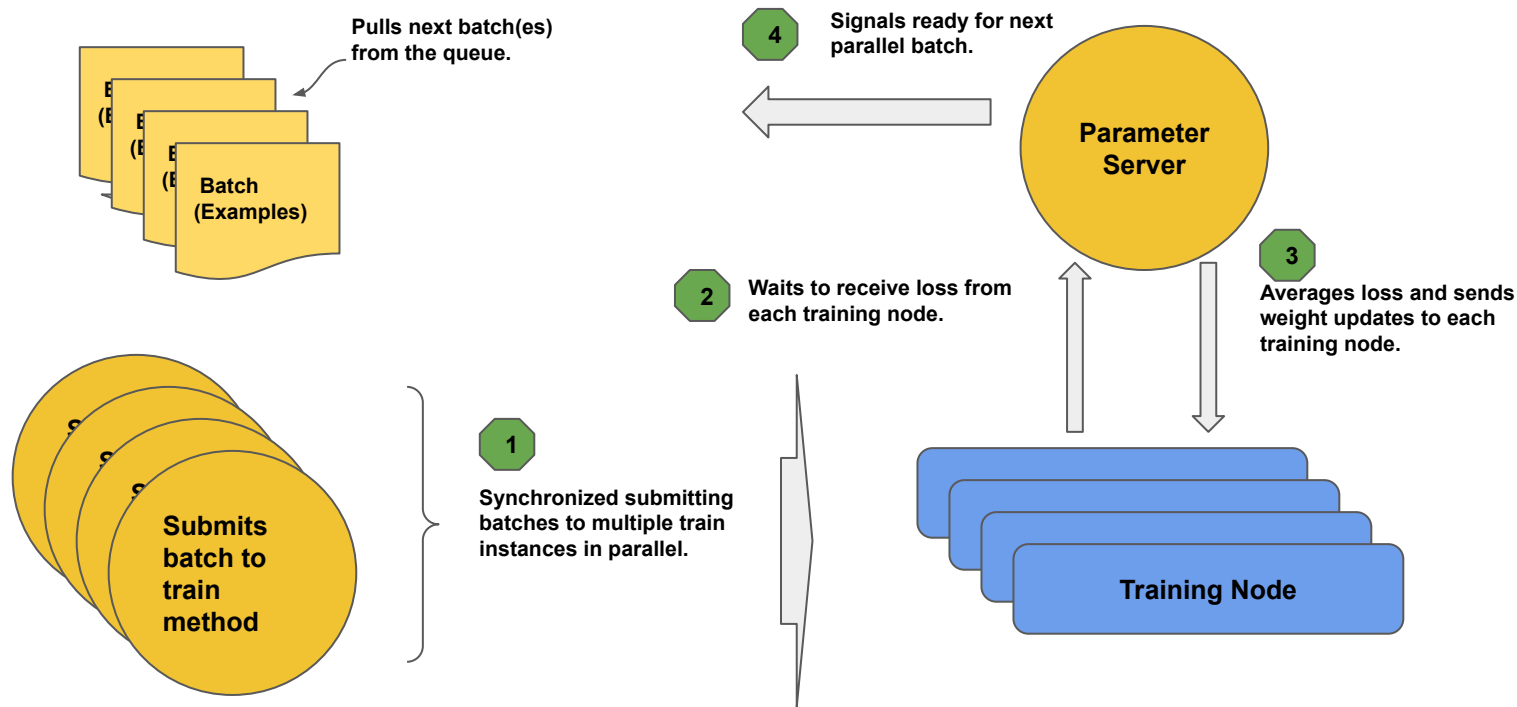
# ML end-2-end production pipeline

- Training Pipeline



# ML end-2-end production pipeline

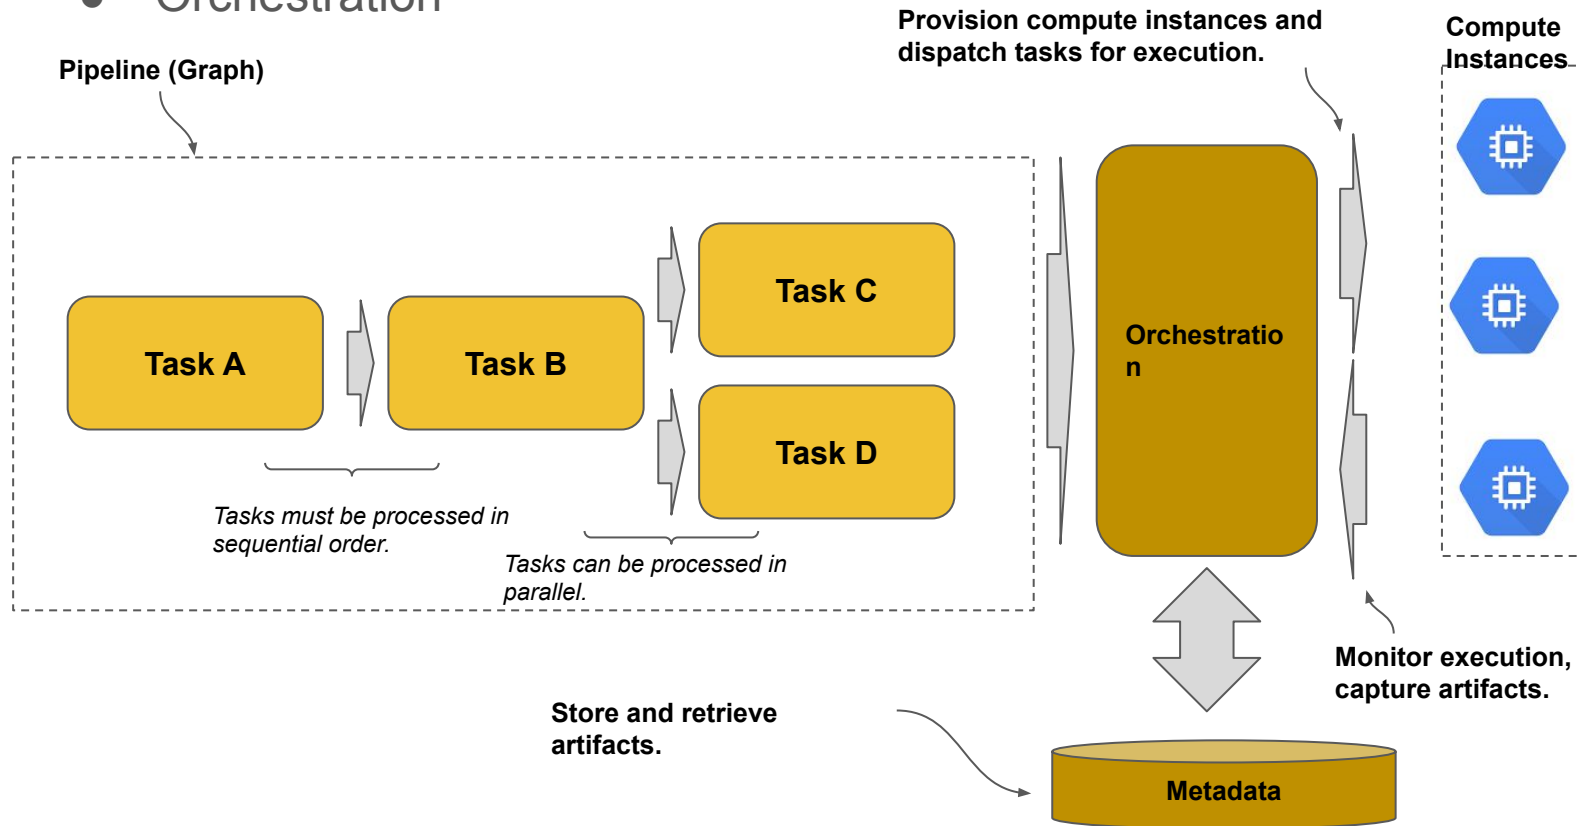
- Training Pipeline





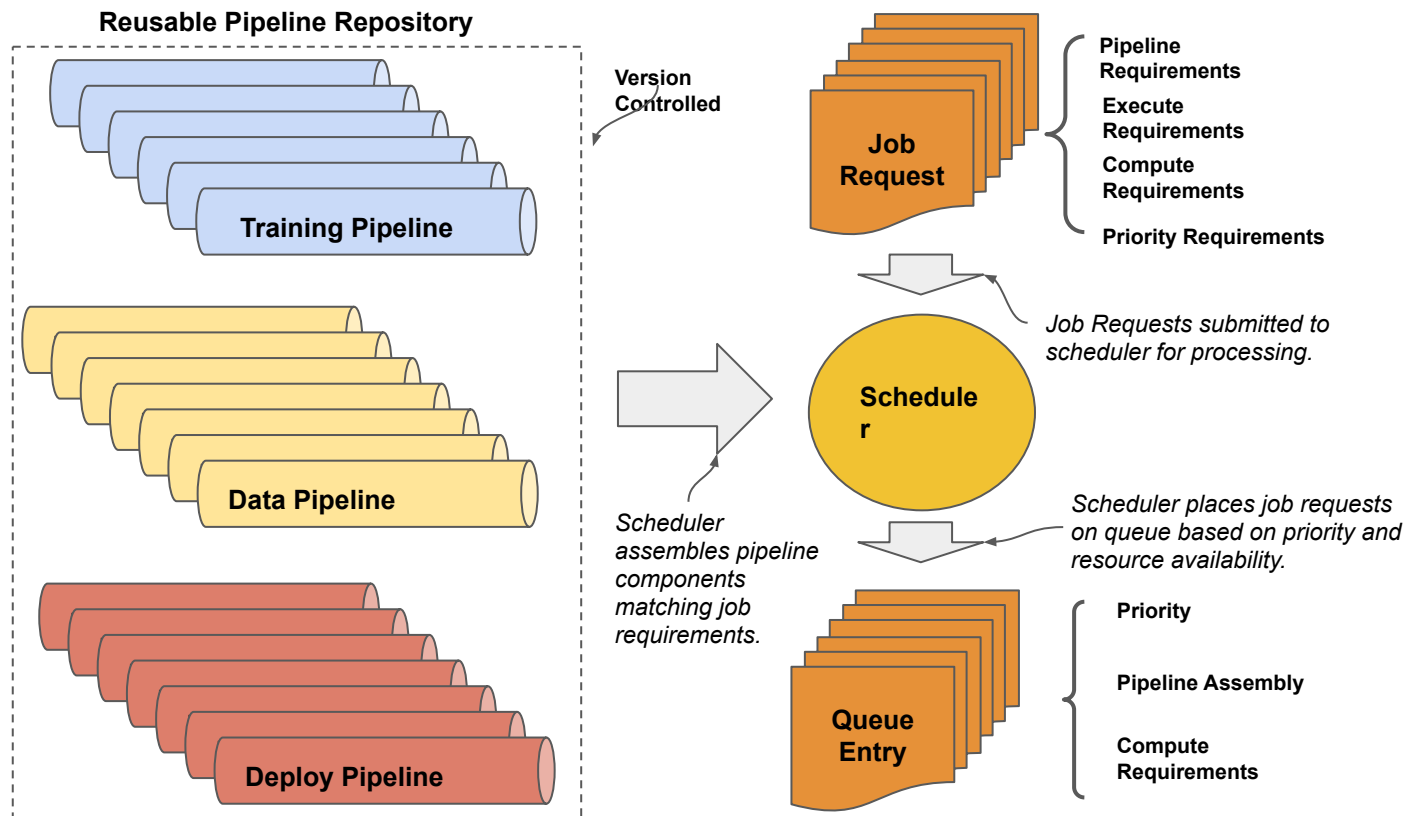
# ML end-2-end production pipeline

- Orchestration



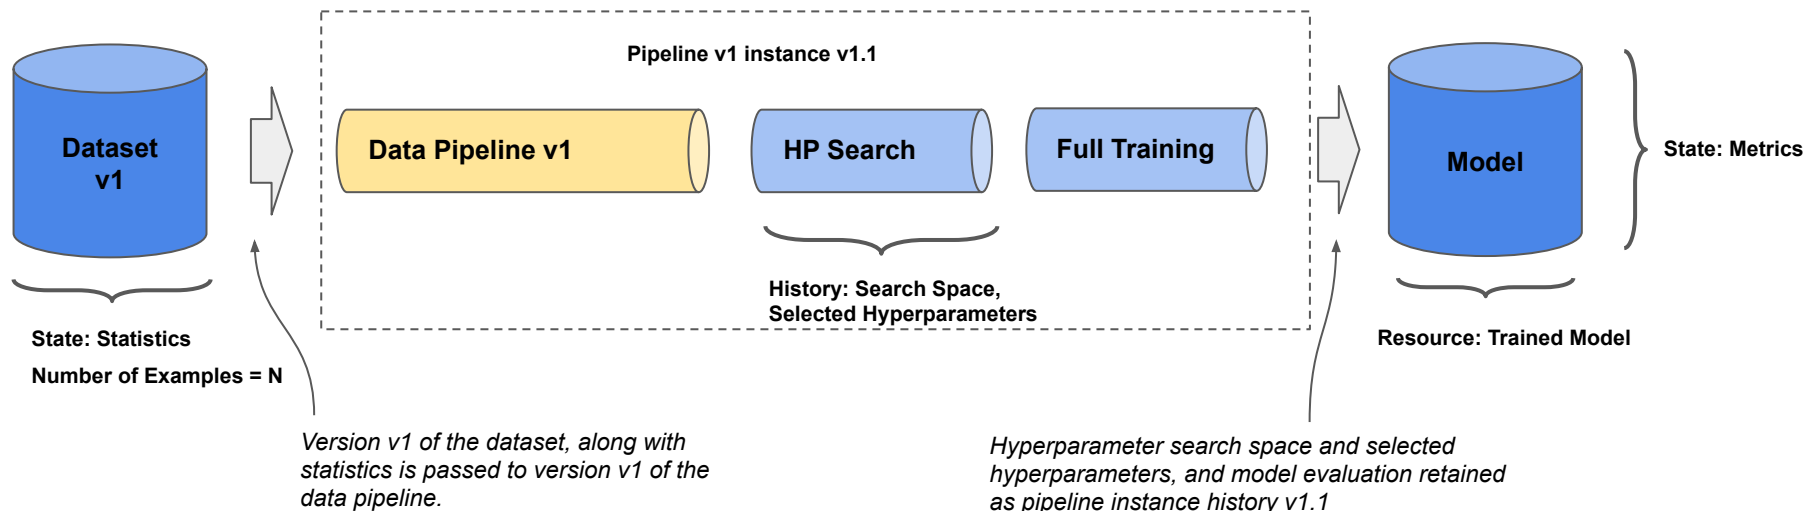
# ML end-2-end production pipeline

- Pipeline Components



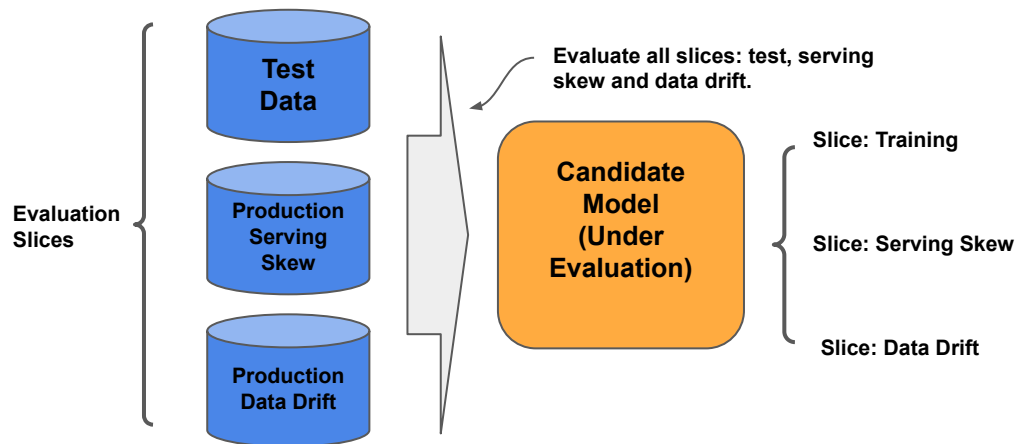
# ML end-2-end production pipeline

- Heuristics



# ML end-2-end production pipeline

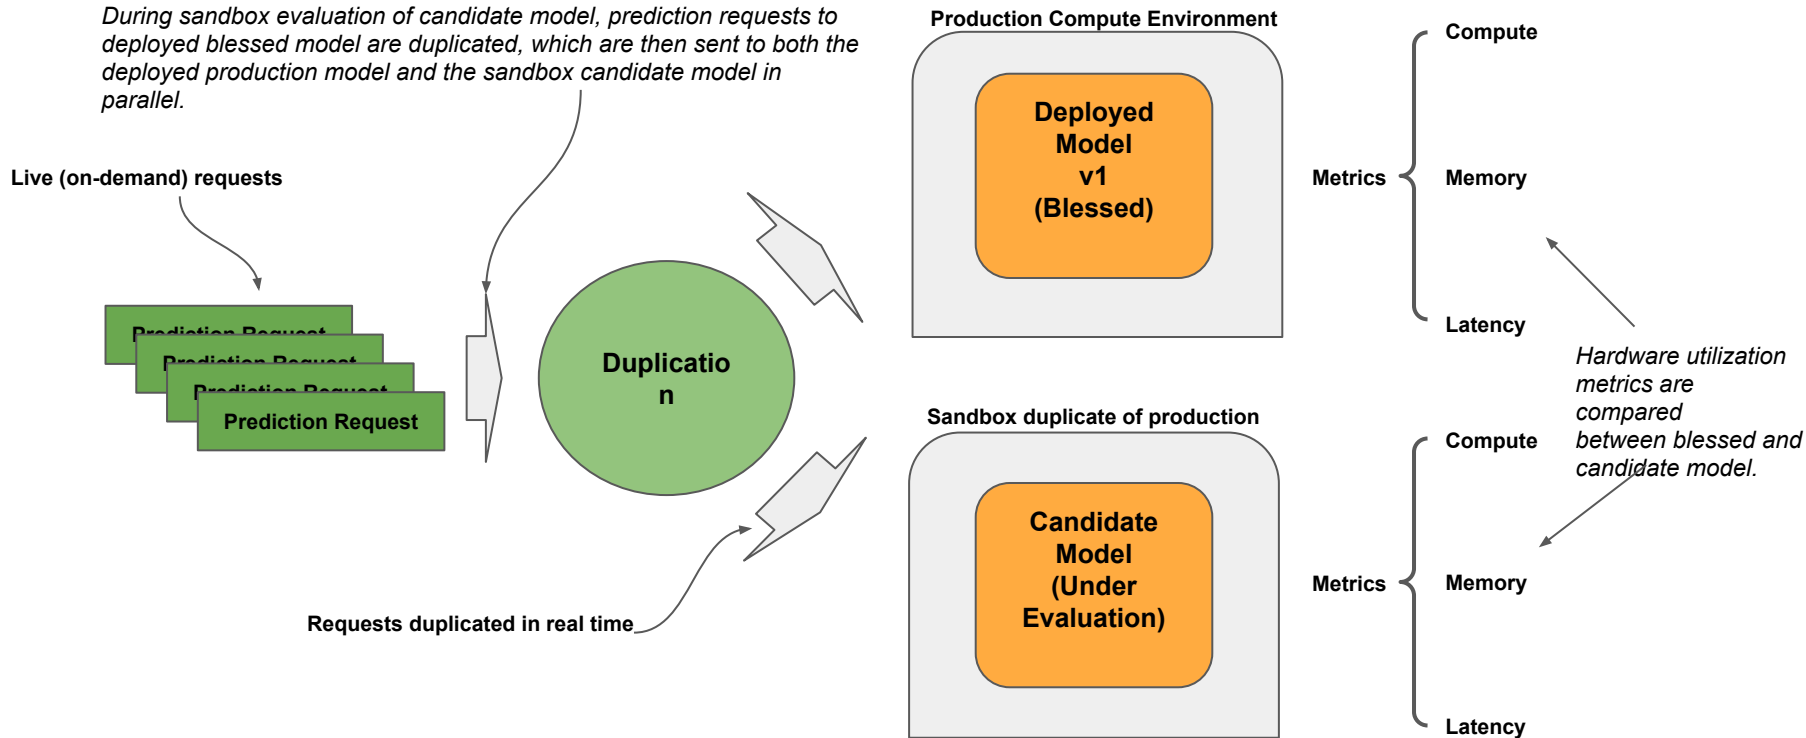
- Evaluation Slicing



# ML end-2-end production pipeline

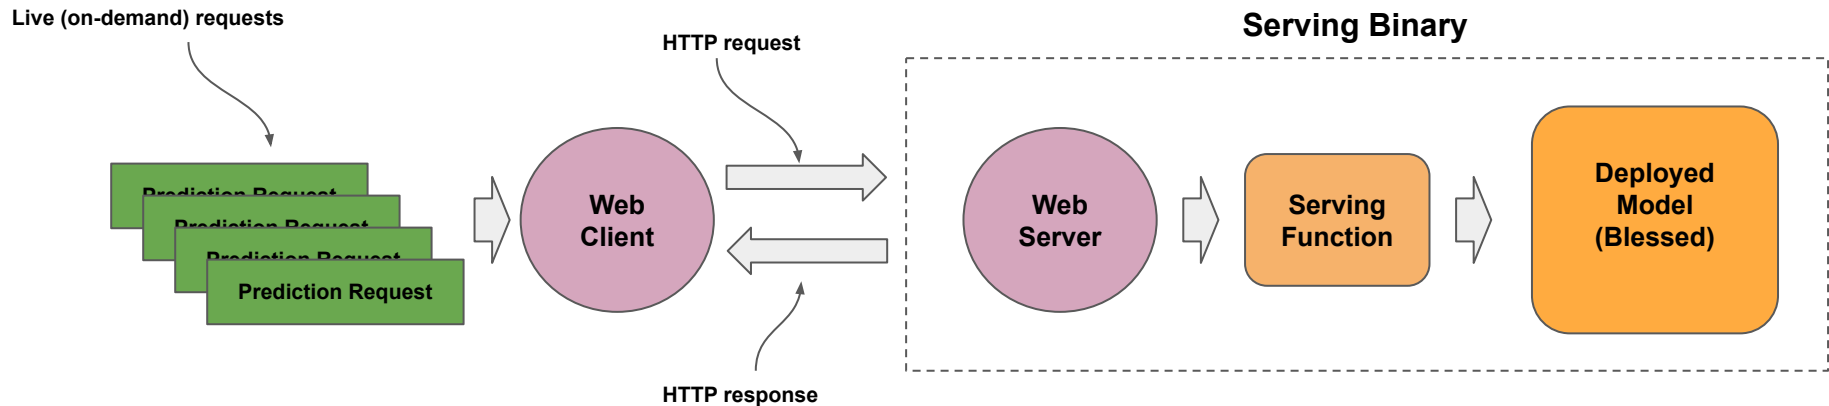
- Sandboxing

*During sandbox evaluation of candidate model, prediction requests to deployed blessed model are duplicated, which are then sent to both the deployed production model and the sandbox candidate model in parallel.*



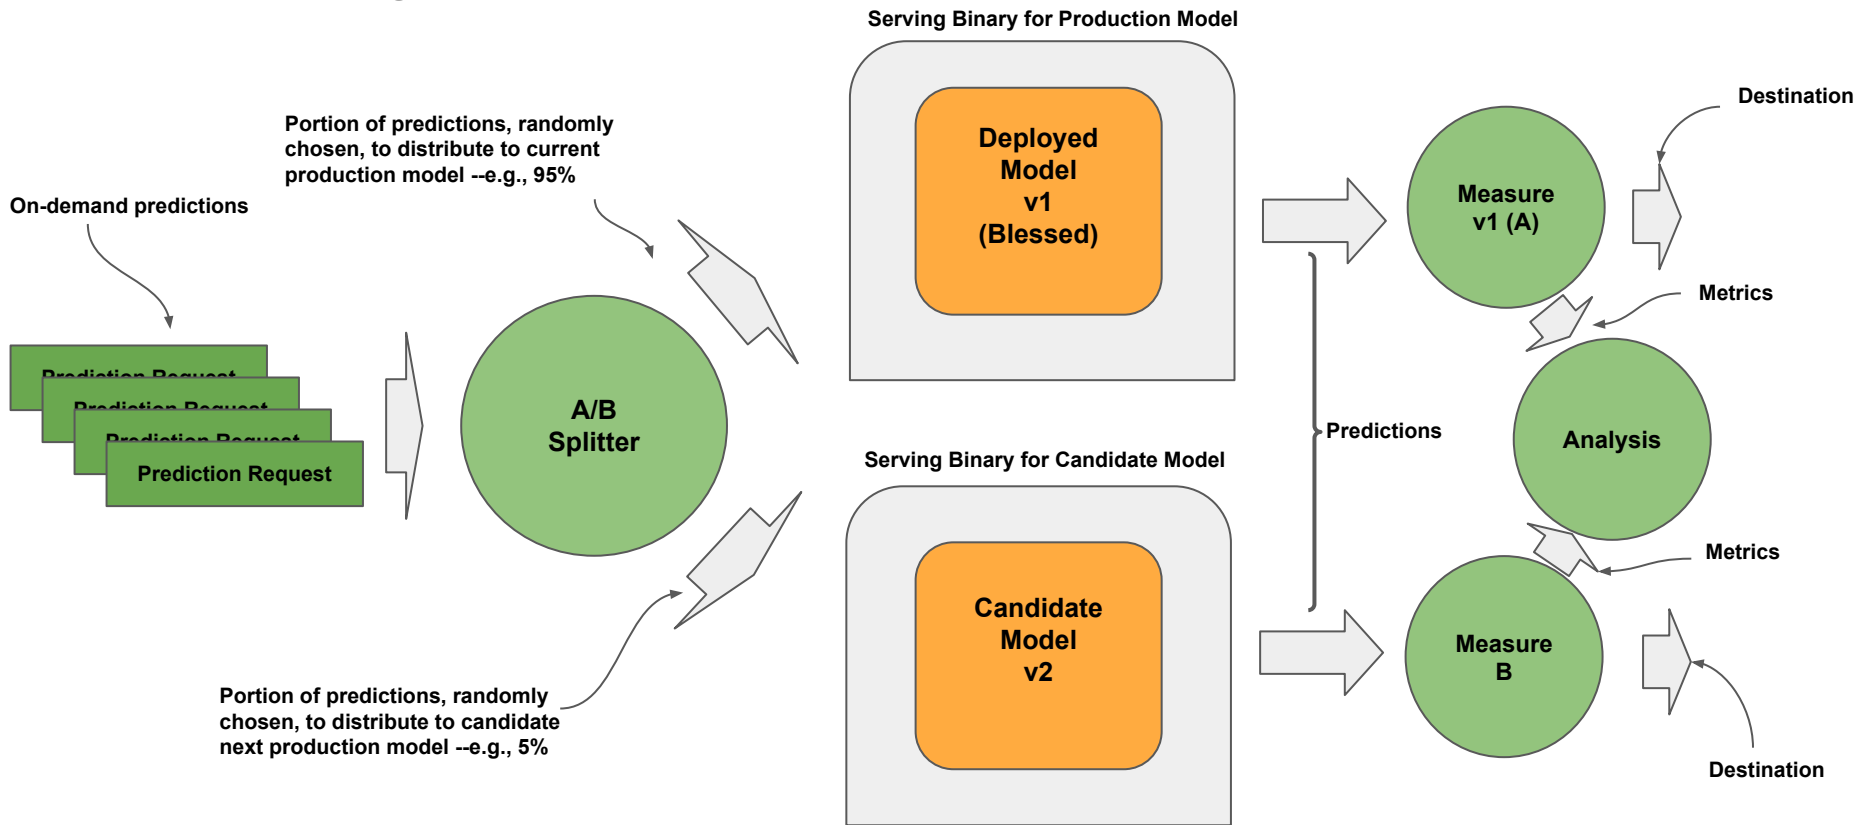
# ML end-2-end production pipeline

- Serving Containers



# ML end-2-end production pipeline

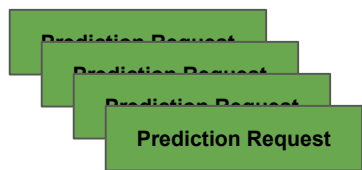
- A/B Testing



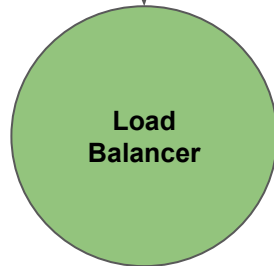
# ML end-2-end production pipeline

- Load Balancing

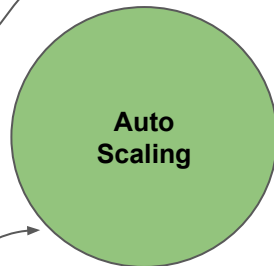
On-demand predictions



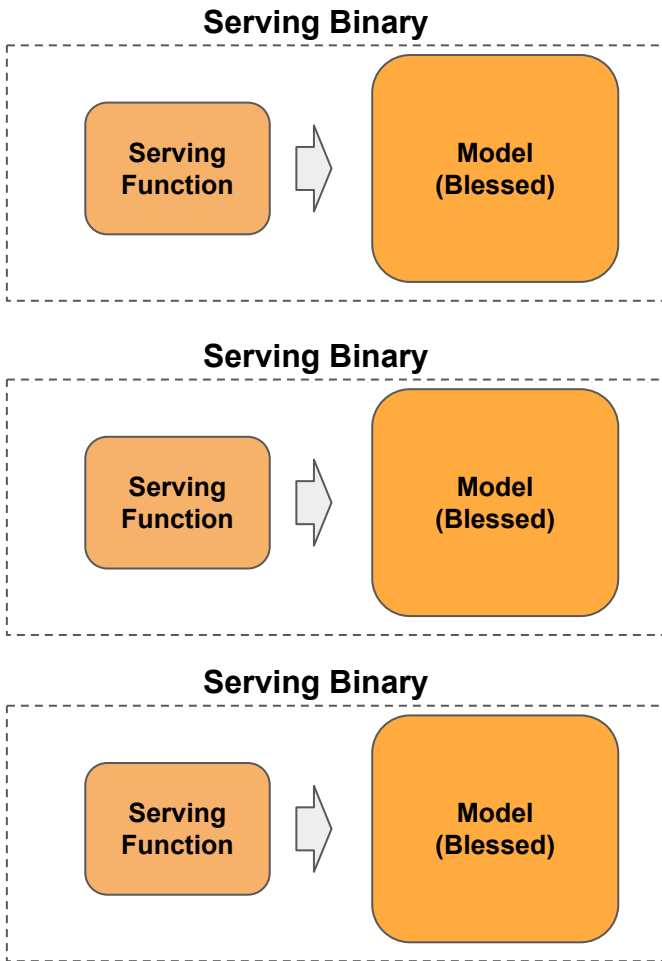
Distributes requests across serving binaries.



Request Frequency,  
Response Latency



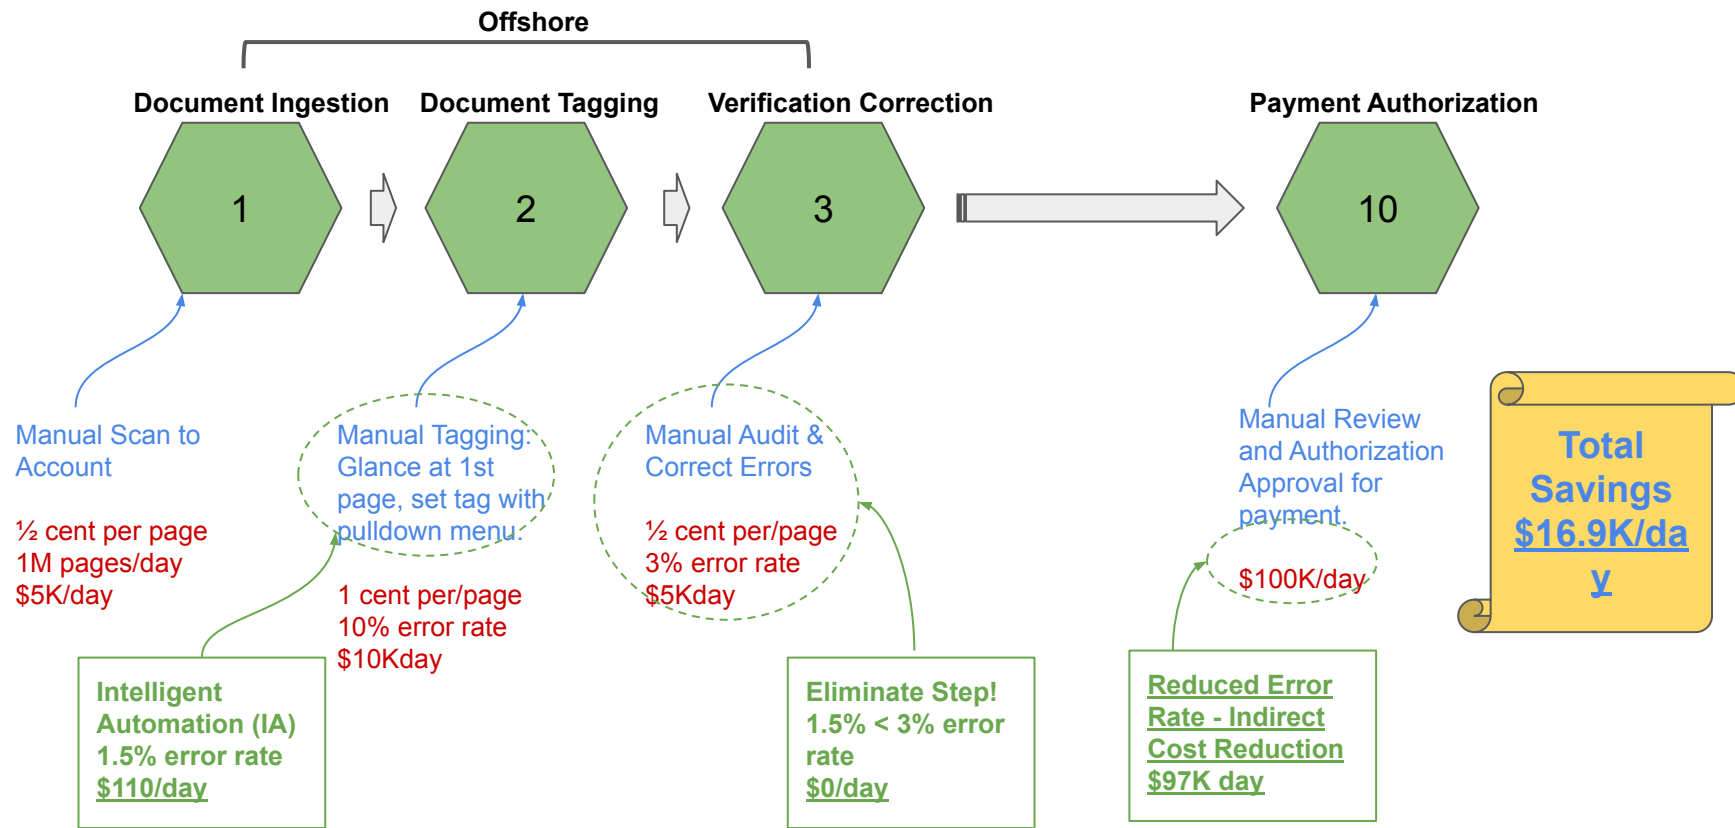
Auto provision and deprovision  
(scaling) serving binary  
instances.





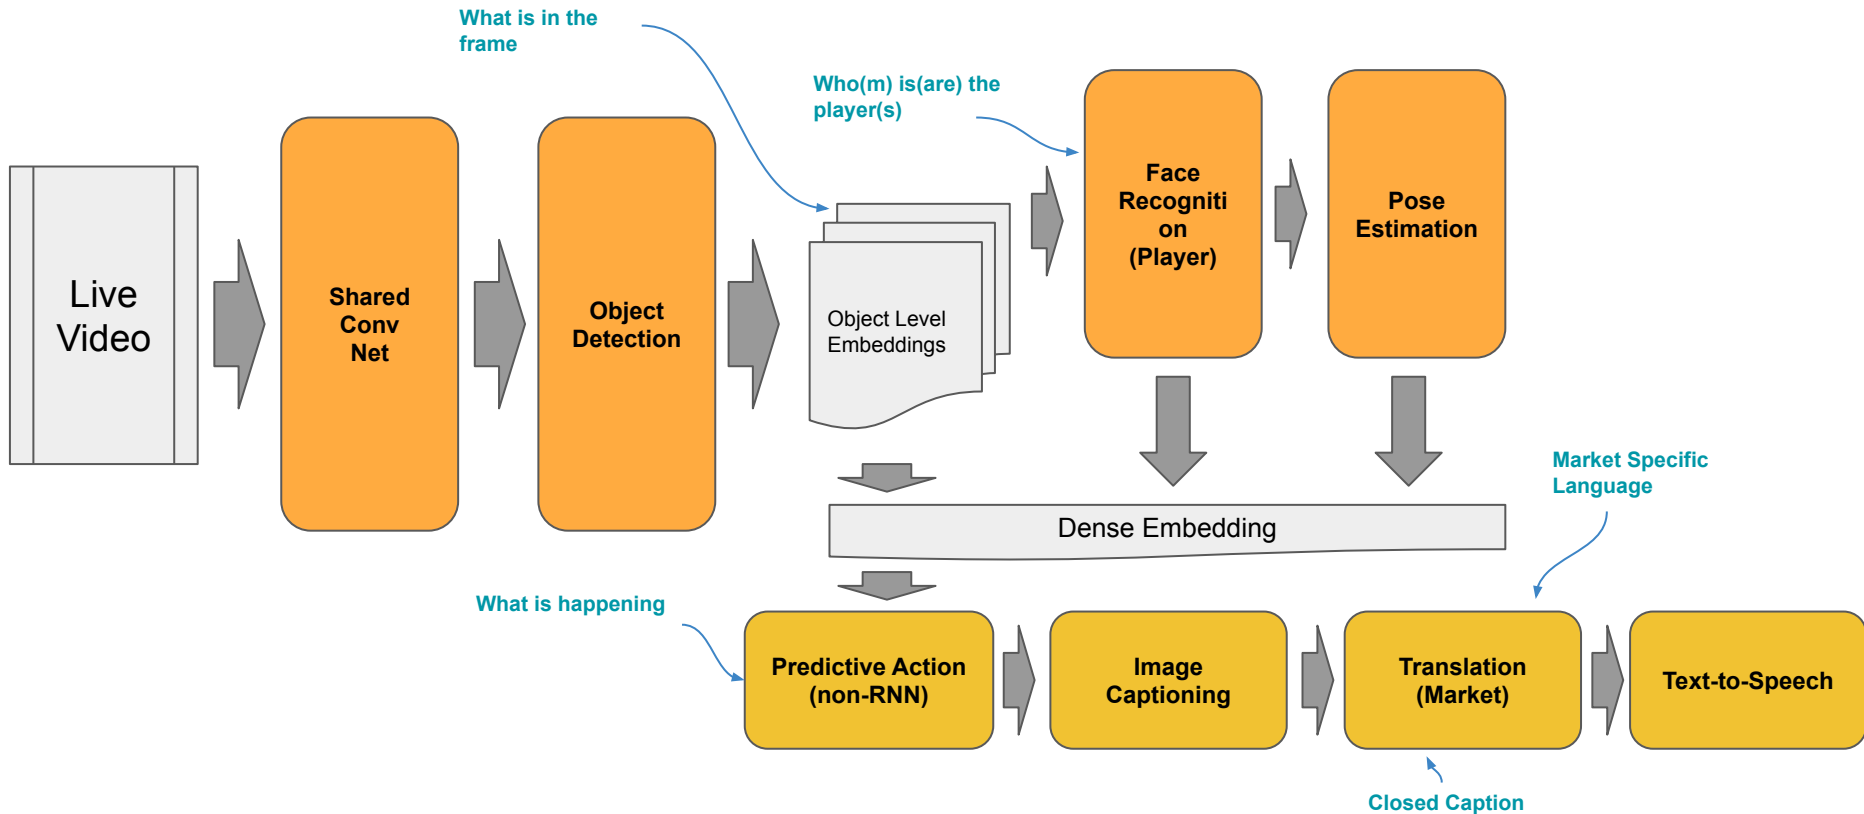
# Framing a Business Problem into an e2e Pipeline

## Intelligent Automation (IA) Applied to Claim Processing



# Framing a Business Problem into an e2e Pipeline

## Model Amalgamation Sports Broadcasting



# AI Platform (Unified) documentation

Let's visit the official documentation for AI Platform (Unified).

AI Platform (Unified) has the following interfaces:

- User Interface
- Command Line (gcloud)
- REST
- **Client Library (SDK)**

The link below takes you to the home page:

<https://cloud.google.com/ai-platform-unified/docs/start/introduction-unified-platform>

# AI Platform (Unified) walk thru

Let's now to the AI Platform (Unified) dashboard (UI). I will walk you through:

- Menu options and selections.
- Creating notebook instances.
- Start/Stop/Open notebook instance.
- Resources: Dataset, Model, Endpoint, Pipeline, etc

# Reducing Costs

- Notebook Instance
  - You don't need a GPU for this training course, so don't select (pay) for one.
  - Select standard instance: 4 vCPUs, 15 GB RAM
  - You pay for each hour the instance is running.
    - 14 cents/hour, ~\$3.36/day
  - Shutdown the instance when not using it (from UI console).

# Reducing Costs

- Deployed Models

- You pay for each hour a model is deployed.
- Deploy the model to the lowest HW configuration
  - single node, n1-standard-4, CPU
- After an exercise, undeploy the model (optionally from UI console).
- Custom Models
  - 19 cents/hour, ~\$4.50/day
- AutoML Models are more pricey
  - image classification: \$1.25/hour, \$30/day
  - object detection: \$1.82/hour, \$44/day
  - Text models: 5 cents/hour, \$1.20/day
  - Tabular models: same as custom, \$4.50/day
- Deployed models get billed a minimum of one hour

# Reducing Costs

- Training
  - AutoML Training
    - Image models: \$3.15/hour
    - Text models: \$3.00/hour
    - Tabular models: \$19/hour
    - Video models: \$2.94/hour
    - Edge models
      - Classification: \$5/hour
      - Object Detection: \$18/hour
    - Use very small size datasets
  - Custom Training
    - 19 cents/hour
    - Do only a few epochs
- <https://cloud.google.com/ai-platform-unified/pricing>

# Reducing Costs

- Strategy for workshop notebooks
  - AutoML
    - follow along (execute) upto training
    - From training on, read only
  - Custom Jobs
    - Execute entire notebook

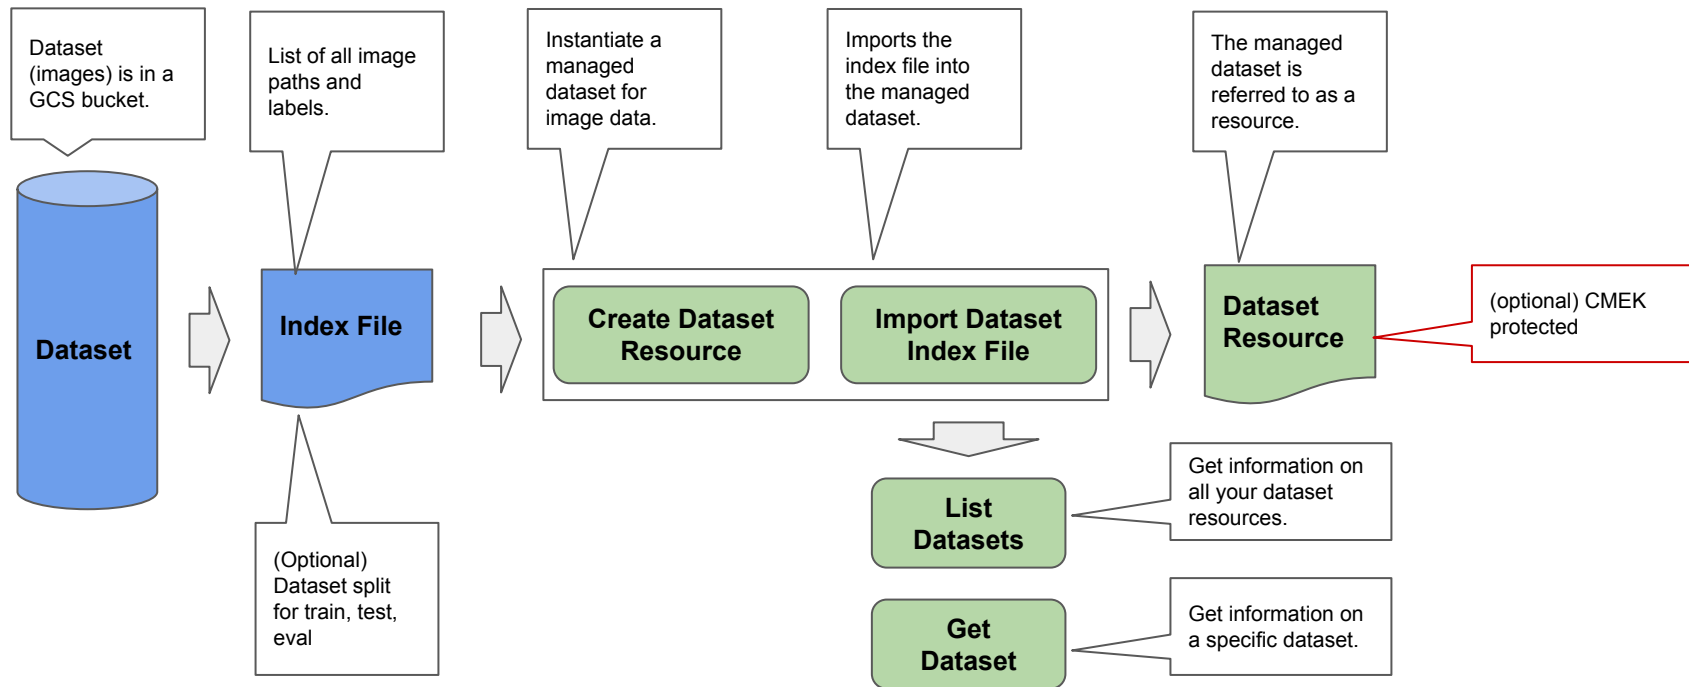


# Workshop 1: AutoML Image Classification

- Create a dataset
- Train a model
- Evaluate the model
- Deploy the model for serving
- Do online prediction

# Workshop 1: AutoML Image Classification

## Create a Dataset



## Create Dataset Resource

### Step 1:

- Instantiate a Dataset resource
- Specify schema for data type
- Optionally user-defined metadata.

### Step 2:

- Create an instance of the Dataset resource.

### Step 3:

- Wait for instance to be created, ~15secs

```
def create_dataset(name, schema, labels=None, timeout=TIMEOUT):
    start_time = time.time()
    try:
        dataset = aip.Dataset(display_name=name,
                               metadata_schema_uri="gs://" + schema,
                               labels=labels)

        operation = clients['dataset'].create_dataset(parent=PARENT, dataset=dataset)

        print("Long running operation:", operation.operation.name)
        result = operation.result(timeout=TIMEOUT)
        print("time:", time.time() - start_time)
        print("response")
        print(" name:", result.name)
        print(" display_name:", result.display_name)
        print(" metadata_schema_uri:", result.metadata_schema_uri)
        print(" metadata:", dict(result.metadata))
        print(" create_time:", result.create_time)
        print(" update_time:", result.update_time)
        print(" etag:", result.etag)
        print(" labels:", dict(result.labels))
        return result
    except Exception as e:
        print("exception:", e)
        return None
```

```
result = create_dataset("flowers-" + TIMESTAMP, DATA_SCHEMA)
```

## Import Dataset Index File

### Step 1:

- Set data labeling schema
- Specify one or more index files.

### Step 2:

- Import the data.

### Step 3:

- Wait for import to complete. Typically a few minutes.

```
def import_data(dataset, gcs_sources, schema):
    config = [{
        'gcs_source': {'uris': gcs_sources},
        'import_schema_uri': schema
    }]

    print("dataset:", dataset_id)
    start_time = time.time()
    try:
        operation = clients['dataset'].import_data(name=dataset_id,
            import_configs=config)
        print("Long running operation:", operation.operation.name)

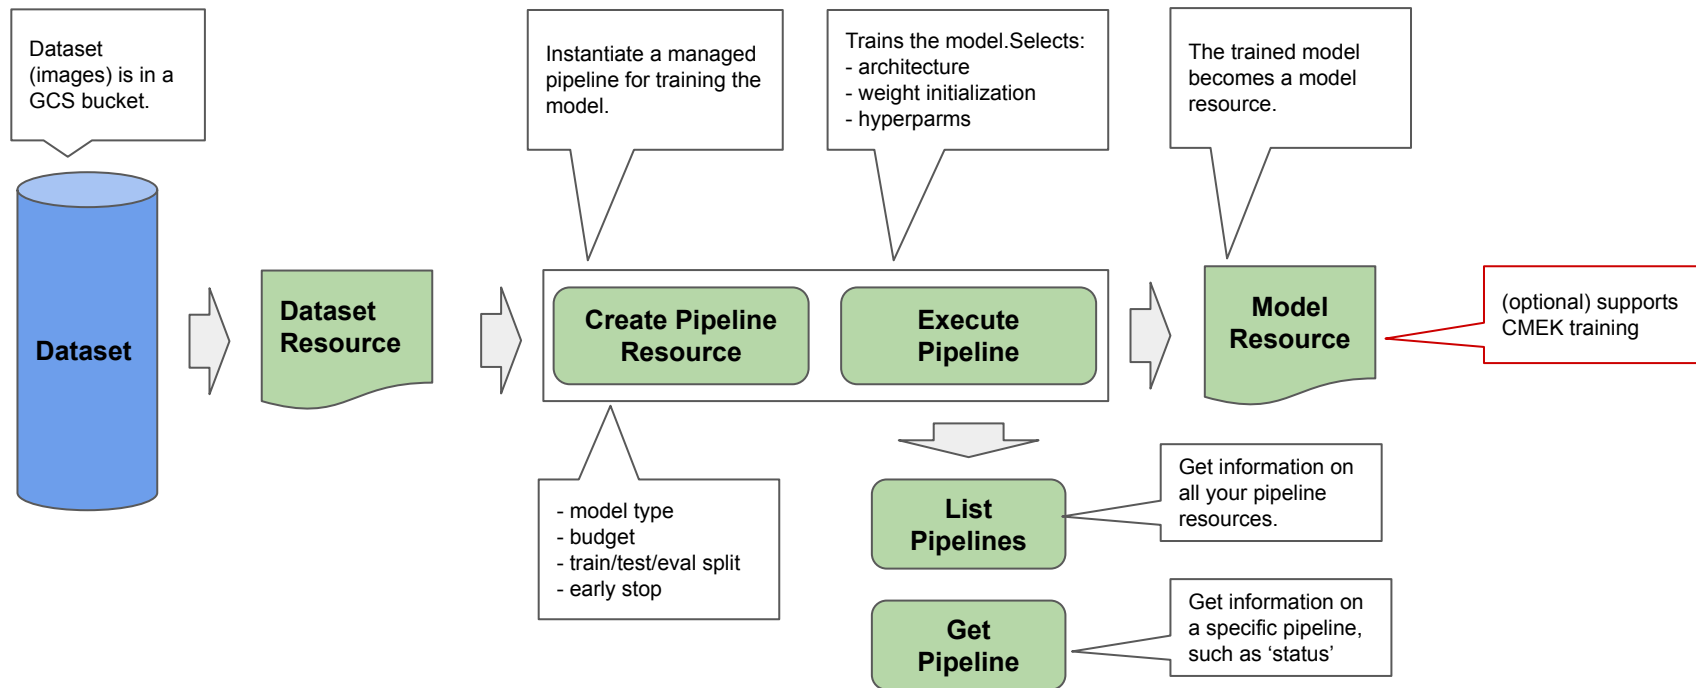
        result = operation.result()
        print("result:", result)
        print("time:", int(time.time() - start_time), "secs")
        print("error:", operation.exception())
        print("meta :", operation.metadata)
        print("after: running:", operation.running(),
            "done:", operation.done(),
            "cancelled:", operation.cancelled())

    return operation
except Exception as e:
    print("exception:", e)
    return None
```

```
import_data(dataset_id, [IMPORT_FILE], LABEL_SCHEMA)
```

# Workshop 1: AutoML Image Classification

## Train a Model



## Create Pipeline Resource

Step 1: Specify the training data input

- Specify the dataset
- Specify the training split.

Step 2: Specify the training pipeline.

- Specify training schema
- Specify task requirements
- Specify training data input
- Human readable name for pipeline and uploaded model.

Step 3:

- Start the training ~ asynchronous

```
def create_pipeline(pipeline_name, model_name, dataset, schema, task):  
  
    dataset_id = dataset.split('/')[1]  
  
    input_config = {'dataset_id': dataset_id,  
                    'fraction_split': {  
                        'training_fraction': 0.8,  
                        'validation_fraction': 0.1,  
                        'test_fraction': 0.1  
                    }}  
  
    training_pipeline = {  
        "display_name": pipeline_name,  
        "training_task_definition": schema,  
        "training_task_inputs": task,  
        "input_data_config": input_config,  
        "model_to_upload": {"display_name": model_name},  
    }  
  
    try:  
        pipeline = clients['pipeline'].create_training_pipeline(parent=PARENT,  
                                                                training_pipeline=training_pipeline)  
  
        print(pipeline)  
    except Exception as e:  
        print("exception:", e)  
        return None  
    return pipeline
```

## Execute Pipeline

Step 1: Query for the training job status.

Step 2: return the status

Step 3: Check for status completion. Will automatically deploy trained model to endpoint for serving

```
def get_training_pipeline(name, silent=False):
    response = clients['pipeline'].get_training_pipeline(name=name)
    if silent:
        return response

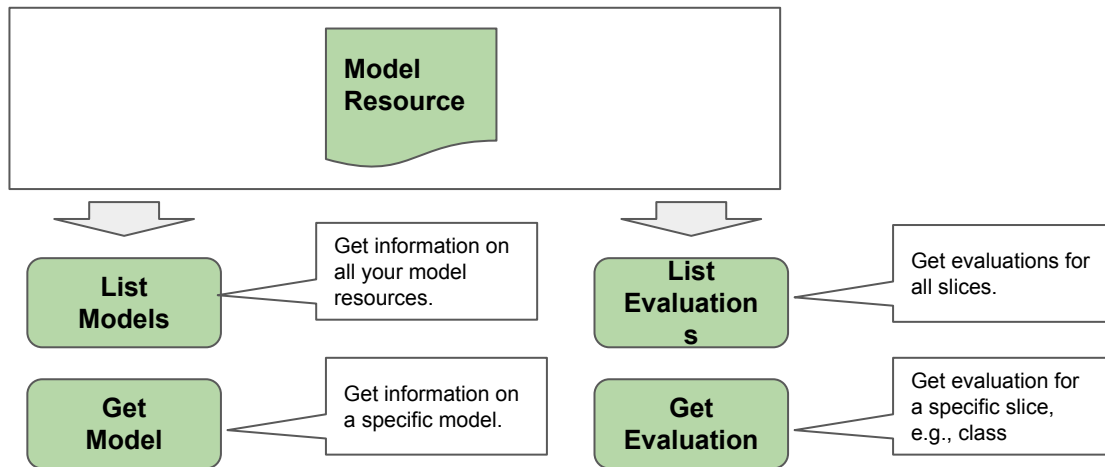
    print("pipeline")
    print(" name:", response.name)
    print(" display_name:", response.display_name)
    print(" state:", response.state)
    print(" training_task_definition:", response.training_task_definition)
    print(" training_task_inputs:", dict(response.training_task_inputs))
    print(" create_time:", response.create_time)
    print(" start_time:", response.start_time)
    print(" end_time:", response.end_time)
    print(" update_time:", response.update_time)
    print(" labels:", dict(response.labels))
    return response

while True:
    response = get_training_pipeline(pipeline_id, True)
    if response.state != aip.PipelineState.PIPELINE_STATE_SUCCEEDED:
        print("Training job has not completed:", response.state)
        model_to_deploy_id = None
        if response.state == aip.PipelineState.PIPELINE_STATE_FAILED:
            raise Exception("Training Job Failed")
        else:
            model_to_deploy = response.model_to_upload
            model_to_deploy_id = model_to_deploy.name
            print("Training Time:", response.end_time - response.start_time)
            break
    time.sleep(60)

print("model to deploy:", model_to_deploy_id)
```

# Workshop 1: AutoML Image Classification

## Evaluate the Model





## List Models

## Get Model

Step 1: Query for information on all trained models (AutoML and Custom)

Step 2: Iterate through the list of model information.

Step 3: Get information on a specific model.

```
def list_models():  
    response = clients['model'].list_models(parent=PARENT)  
    for model in response:  
        print("name", model.name)  
        print("display_name", model.display_name)  
        print("create_time", model.create_time)  
        print("update_time", model.update_time)  
        print("container", model.container_spec.image_uri)  
        print("artifact_uri", model.artifact_uri)  
        print('\n')  
    return response
```

list\_models()

```
def get_model(name):  
    response = clients['model'].get_model(name=name)  
    print(response)
```

get\_model(model\_to\_deploy\_name)

## List Evaluations

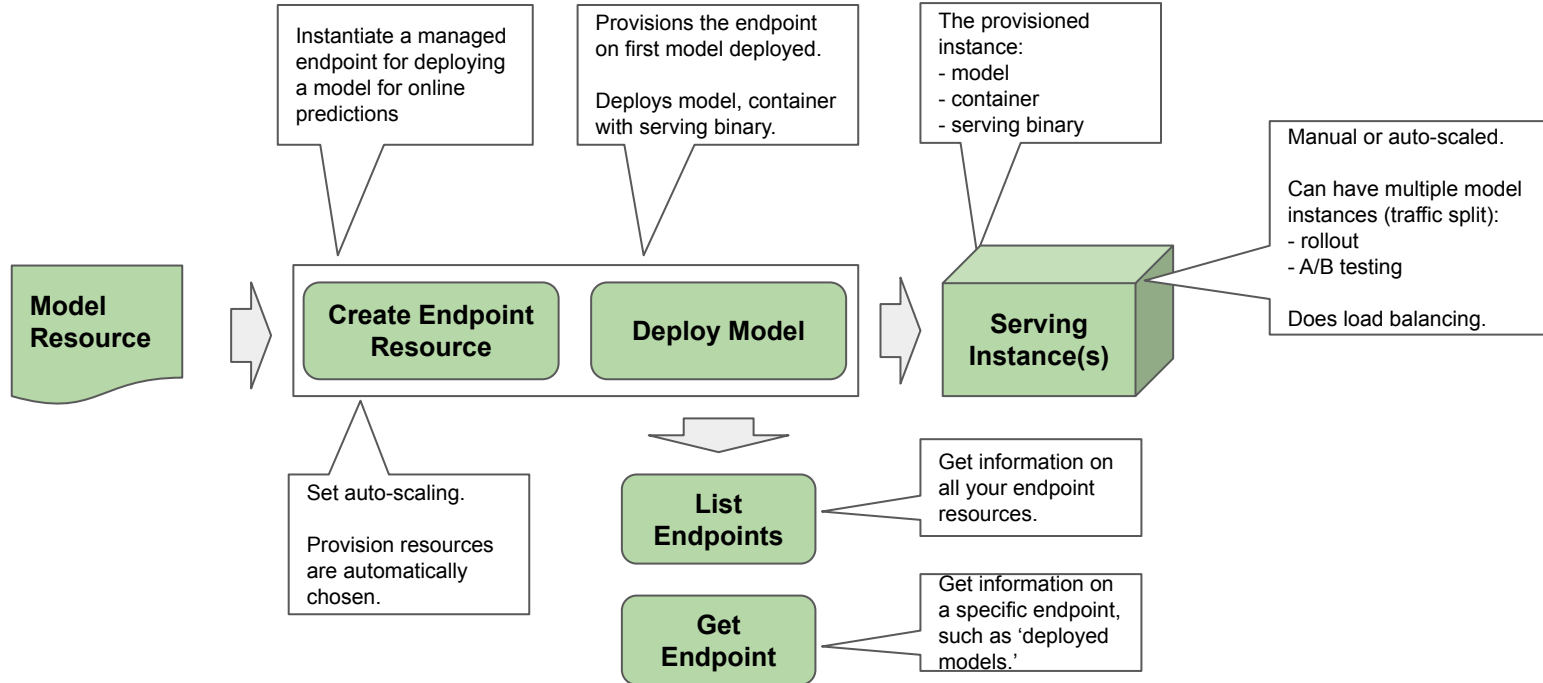
Step 1: Query for evaluations on all slices of the test/eval data (e.g., by class)

Step 2: Iterate through the list of evaluation slices.

```
def list_model_evaluations(name):  
    response = clients['model'].list_model_evaluations(parent=name)  
    for evaluation in response:  
        print("model_evaluation")  
        print(" name:", evaluation.name)  
        print(" metrics_schema_uri:", evaluation.metrics_schema_uri)  
        metrics = json_format.MessageToDict(evaluation._pb.metrics)  
        for metric in metrics.keys():  
            print(metric)  
            print('logloss', metrics['logLoss'])  
            print('auPrc', metrics['auPrc'])  
  
    return response  
  
list_model_evaluations(model_to_deploy_id)
```

# Workshop 1: AutoML Image Classification

## Deploy for Serving



## Create Endpoint Resource

Step 1: Create Endpoint resource.  
Automatically chooses HW for deployment.

Step 2: Wait for endpoint to be created.

Step 3: Get the endpoint ID

```
def create_endpoint(display_name):
    endpoint = {"display_name": display_name}
    response = clients['endpoint'].create_endpoint(parent=PARENT,
                                                    endpoint=endpoint)
    print("Long running operation:", response.operation.name)

    result = response.result(timeout=300)
    print("result")
    print(" name:", result.name)
    print(" display_name:", result.display_name)
    print(" description:", result.description)
    print(" labels:", result.labels)
    print(" create_time:", result.create_time)
    print(" update_time:", result.update_time)
    return result
```

## Deploy Model

Step 1: Specify the model to deploy, and manual/auto-scaling settings.

Step 2:  
- Specify the traffic split  
- Deploy the model

Step 3:  
- Wait for model deployed to complete.

```
def deploy_model(model, deployed_model_display_name, endpoint,
                 traffic_split={"0": 100}):

    deployed_model = {
        "model": model,
        "display_name": deployed_model_display_name,
        "automatic_resources": {
            "min_replica_count": MIN_NODES,
            "max_replica_count": MAX_NODES
        },
    },
}

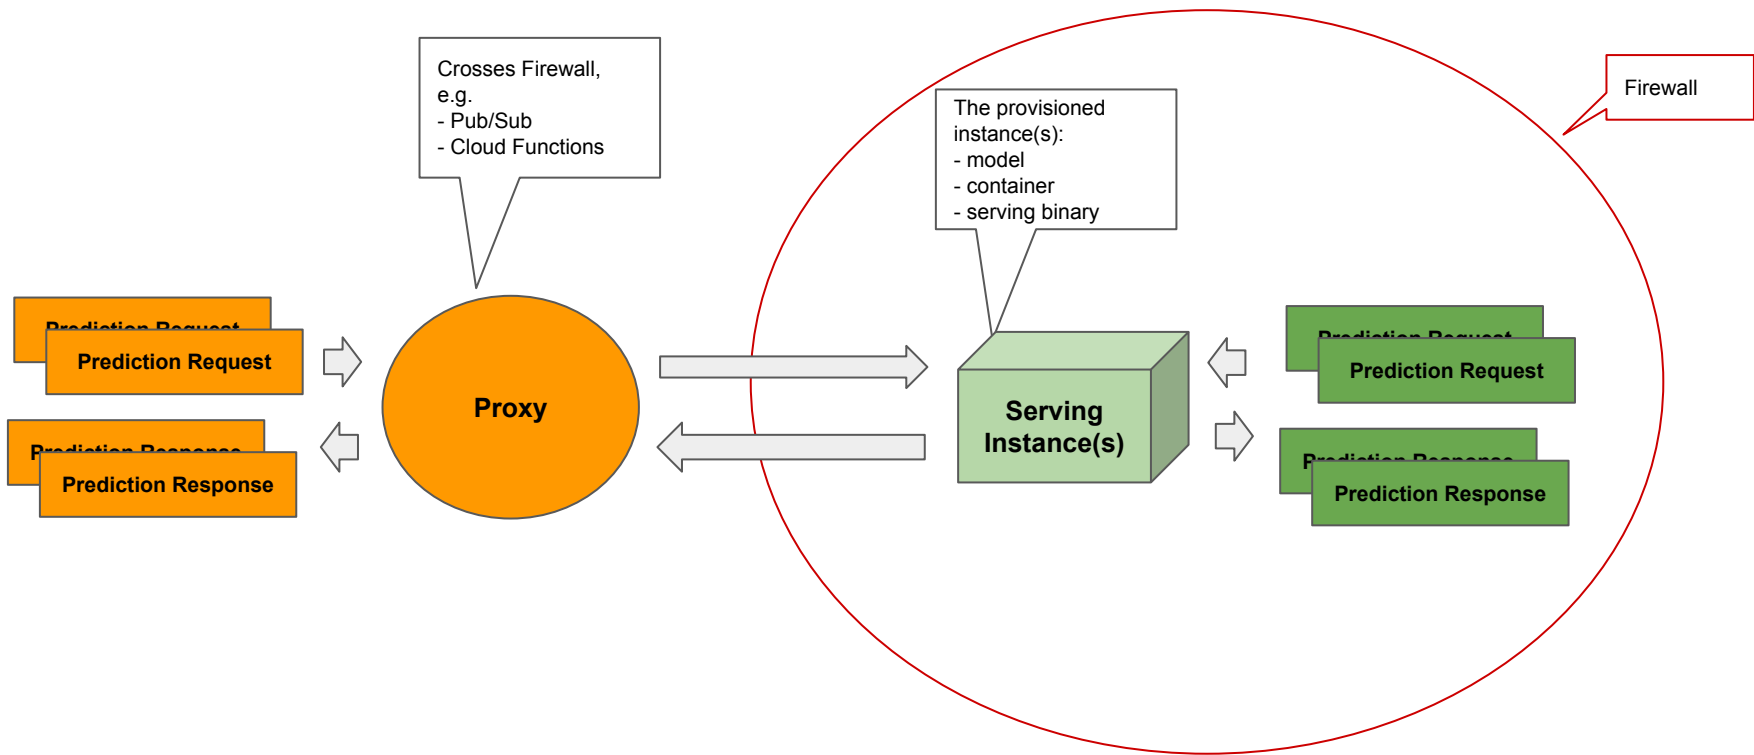
    response = clients['endpoint'].deploy_model(
        endpoint=endpoint, deployed_model=deployed_model, traffic_split=traffic_split)

    print("Long running operation:", response.operation.name)
    result = response.result()
    print("result")
    deployed_model = result.deployed_model
    print(" deployed_model")
    print(" id:", deployed_model.id)
    print(" model:", deployed_model.model)
    print(" display_name:", deployed_model.display_name)
    print(" create_time:", deployed_model.create_time)

    return deployed_model.id
```

# Workshop 1: AutoML Image Classification

## Do Online Predictions



## Serving

Step 1: Get compressed image bytes

Step 2:  
- base64 encode the image

Step 3:  
- Construct list of instances to predict.

Step 4:  
- Make prediction request  
- Set parameters for returning results.

```
def predict_item(filename, endpoint, parameters_dict):  
    parameters = json_format.ParseDict(parameters_dict, Value())  
  
    with tf.io.gfile.GFile(filename, "rb") as f:  
        content = f.read()  
  
    instances_list = [{"content": base64.b64encode(content).decode("utf-8")}]  
    instances = [json_format.ParseDict(s, Value()) for s in instances_list]  
  
    response = clients['prediction'].predict(endpoint=endpoint, instances=instances,  
        parameters=parameters)  
    print("response")  
    print(" deployed_model_id:", response.deployed_model_id)  
    predictions = response.predictions  
    print("predictions")  
    for prediction in predictions:  
        print(" prediction:", dict(prediction))  
  
predict_item(test_item, endpoint_id,  
             {'confidenceThreshold': 0.5, 'maxPredictions': 2})
```

## Workshop 2: AutoML Image Batch, IOD, ISG, Edge

- Create a batch job for image classification
- Train an image object detection model
- Train an image segmentation
- Export a model for Edge prediction
- Do edge prediction



# Workshop 2: AutoML Batch Prediction

## Make Batch File

Step 1: Set paths to the images stored in GCS

Step 2: Create JSONL file on GCS

Step 3: Write each instance to predict as a JSON object.

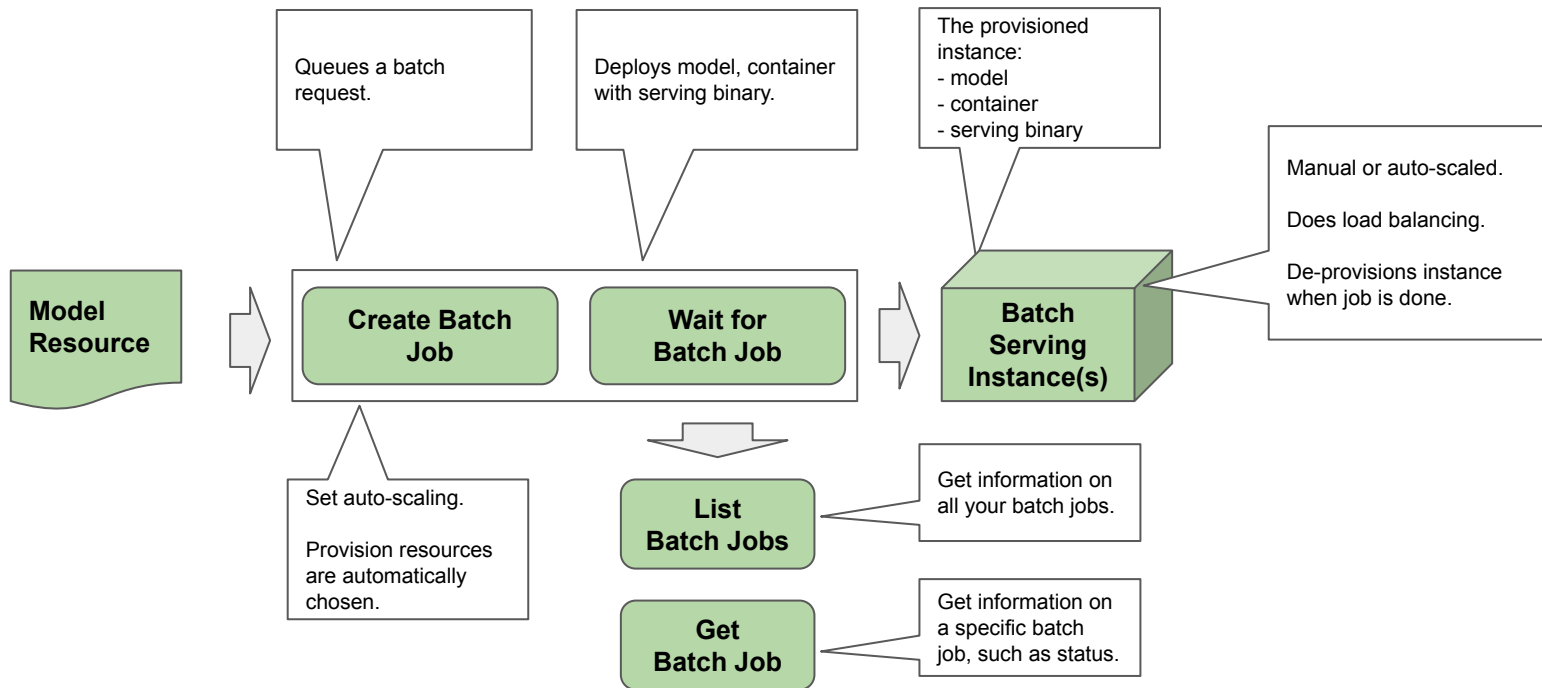
```
test_item_1 = BUCKET_NAME + "/" + file_1
test_item_2 = BUCKET_NAME + "/" + file_2

import tensorflow as tf
import json

gcs_input_uri = BUCKET_NAME + '/test.jsonl'
with tf.io.gfile.GFile(gcs_input_uri, 'w') as f:
    data = {"content": test_item_1, "mime_type": "image/jpeg"}
    f.write(json.dumps(data) + '\n')
    data = {"content": test_item_2, "mime_type": "image/jpeg"}
    f.write(json.dumps(data) + '\n')
```

# Workshop 2: AutoML Batch Prediction

## Make Batch Request - No Endpoint/Deployed Model



## Create Batch Job

Step 1: Specify HW resources for each VM instance.

Step 2: Create requirements spec for batch job.

Step 3: Specify one or more batch input files as a list.

Step 4: Specify location on GCS to store the predictions

Step 5: Set manual/auto scaling

Step 6: Submit the batch job

```
def create_batch_prediction_job(display_name, model_name, gcs_source_uri,
                               gcs_destination_output_uri_prefix, parameters):
    if DEPLOY_GPU:
        machine_spec = {
            "machine_type": DEPLOY_COMPUTE,
            "accelerator_type": DEPLOY_GPU,
            "accelerator_count": DEPLOY_NGPU,
        }
    else:
        machine_spec = {
            "machine_type": DEPLOY_COMPUTE,
            "accelerator_count": 0,
        }

    batch_prediction_job = {
        "display_name": display_name,
        "model": model_name,
        "model_parameters": json_format.ParseDict(parameters, Value()),
        "input_config": {
            "instances_format": IN_FORMAT,
            "gcs_source": {"uris": [gcs_source_uri]},
        },
        "output_config": {
            "predictions_format": OUT_FORMAT,
            "gcs_destination": {"output_uri_prefix": gcs_destination_output_uri_prefix},
        },
        "dedicated_resources": {
            "machine_spec": machine_spec,
            "starting_replica_count": MIN_NODES,
            "max_replica_count": MAX_NODES
        }
    }
    response = clients['job'].create_batch_prediction_job(
        parent=PARENT, batch_prediction_job=batch_prediction_job
    )
    return response

IN_FORMAT = 'jsonl'
OUT_FORMAT = 'jsonl' # [jsonl]

response = create_batch_prediction_job(BATCH_MODEL, model_to_deploy_id, gcs_input_uri, BUCKET_NAME,
                                       {'confidenceThreshold': 0.5, 'maxPredictions': 2})
```


# Workshop 2: AutoML Image Object Detection

**Train Image Object Detection**

## Image Object Detection (IOD) - Schema

```
# Image Dataset type
DATA_SCHEMA = 'google-cloud-aiplatform/schema/dataset/metadata/image_1.0.0.yaml'
# Image Labeling type
LABEL_SCHEMA = "gs://google-cloud-aiplatform/schema/dataset/ioformat/image_bounding_box_io_format_1.0.0.yaml"
# Image Training task
TRAINING_SCHEMA =
"gs://google-cloud-aiplatform/schema/trainingjob/definition/automl_image_object_detection_1.0.0.yaml"
```

LABEL and TRAINING SCHEMA  
specific to IOD




## Image Object Detection (IOD) - Labeling

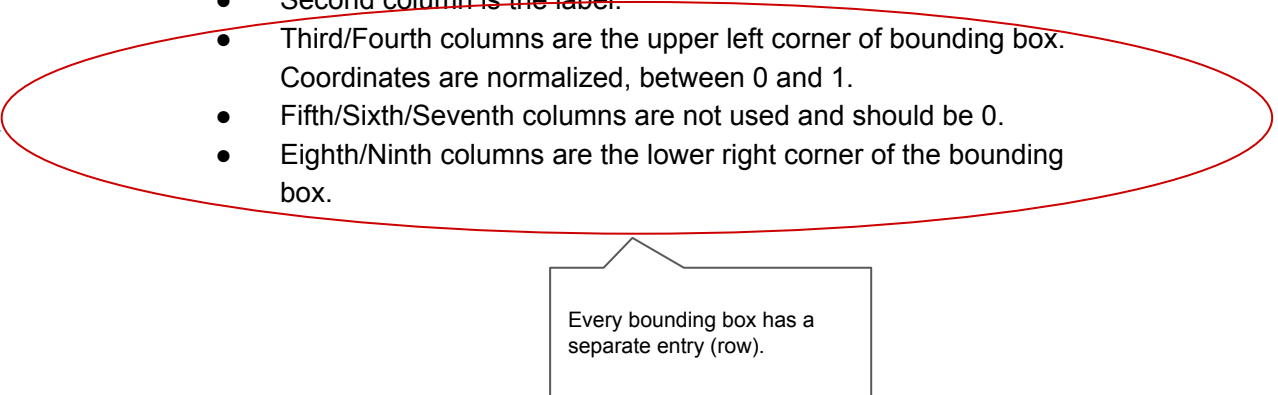
For image object detection, the CSV index file has the requirements:

- No heading.
- First column is the Cloud Storage path to the image.
- ~~Second column is the label.~~
- Third/Fourth columns are the upper left corner of bounding box. Coordinates are normalized, between 0 and 1.
- Fifth/Sixth/Seventh columns are not used and should be 0.
- Eighth/Ninth columns are the lower right corner of the bounding box.

Additional columns for defining the bounding box.



Every bounding box has a separate entry (row).



## Image Object Detection (IOD) - Prediction

The `response` object returns a list, where each element in the list corresponds to the corresponding image in the request. You will see in the output for each prediction:

- - Confidence level in the prediction (confidences).
- - The predicted label (displayNames).
- - The bounding box for the label (bbboxes).

Additional output for the bounding box of each predicted object label.



## Image Object Detection (IOD) - Batch Prediction

For JSONL file, you make one dictionary entry per line for each data item (instance). The dictionary contains the key/value pairs:

- content: The Cloud Storage path to the image.
- mime\_type: The content type. In our example, it is an jpeg file.

For example:

`{'content': '[your-bucket]/file1.jpg', 'mime_type': 'jpeg'}`

Same as image classification





# Workshop 2: AutoML Image Segmentation

## Image Segmentation (ISG) - Schema

```
# Image Dataset type
DATA_SCHEMA = 'google-cloud-aiplatform/schema/dataset/metadata/image_1.0.0.yaml'
# Image Labeling type
LABEL_SCHEMA = "gs://google-cloud-aiplatform/schema/dataset/ioformat/image_segmentation_io_format_1.0.0.yaml"
# Image Training task
TRAINING_SCHEMA =
"gs://google-cloud-aiplatform/schema/trainingjob/definition/automl_image_segmentation_1.0.0.yaml"
```

LABEL and TRAINING SCHEMA  
specific to IOD

## Image Segmentation (ISG) - Labeling

For image segmentation, the JSONL index file has the requirements:

- - Each data item is a separate JSON object, on a separate line.
- - The key/value pair `image\_gcs\_uri` is the Cloud Storage path to the image.
- - The key/value pair `category\_mask\_uri` is the Cloud Storage path to the mask image in PNG format.
- - The key/value pair `annotation\_spec\_colors` is a list mapping mask colors to a label.
- - The key/value pair pair `display\_name` is the label for the pixel color mask.
- - The key/value pair pair `color` are the RGB normalized pixel values (between 0 and 1) of the mask for the corresponding label.

All fields except for image path are specific to segmentation

```
{ 'image_gcs_uri': image, 'segmentation_annotations': { 'category_mask_uri': mask_image, 'annotation_spec_colors' : [ { 'display_name': label, 'color': { "red": value, "blue", value, "green": value } }, ...] }
```

Cleaner to specify as JSON than as CSV.

## Image Segmentation (ISG) - Task Requirements

```
task = json_format.ParseDict({'budget_milli_node_hours': 2000,  
                             'model_type': "CLOUD_LOW_ACCURACY_1"  
                             }, Value())  
  
response = create_pipeline(PIPE_NAME, MODEL_NAME, dataset_id, TRAINING_SCHEMA, task)
```

Model type specific to ISG.


Select either high or low  
accuracy tradeoff for  
size/latency.

## Image Segmentation (ISG) - Prediction

The `response` object returns a list, where each element in the list corresponds to the corresponding image in the request. You will see in the output for each prediction:

- - ConfidenceMask - Confidence level in the prediction
- - CategoryMask - Predictions per pixel.

Output is on a per pixel basis



## Image Segmentation (ISG) - Batch Prediction

For JSONL file, you make one dictionary entry per line for each data item (instance). The dictionary contains the key/value pairs:

- content: The Cloud Storage path to the image.
- mime\_type: The content type. In our example, it is an jpeg file.

For example:


`{'content': '[your-bucket]/file1.jpg', 'mime_type': 'jpeg'}`

Same as image classification



# Workshop 2: AutoML Image Models, Export to Edge

## Deploy for Edge Serving



## Image Model Exported to Edge - Training

```
PIPE_NAME = "salads_pipe-" + TIMESTAMP
MODEL_NAME = "salads_model-" + TIMESTAMP

task = json_format.ParseDict({'budget_milli_node_hours': 20000,
                             'model_type': "MOBILE_TF_LOW_LATENCY_1",
                             'disable_early_stopping': False
                             }, Value())

response = create_pipeline(PIPE_NAME, MODEL_NAME, dataset_id, TRAINING_SCHEMA, task)
```

Model Type are specific to edge models:

- MOBILE\_TF\_LOW\_LATENCY\_1
- MOBILE\_TF\_HIGH\_ACCURACY\_1
- MOBILE\_TF\_VERSATILE\_1

Can train edge model for:

- image classification
- object detection



## Image Model Exported to Edge - Export

```
def export_model(name, format, gcs_dest):  
    output_config = {  
        "artifact_destination": {"output_uri_prefix": gcs_dest},  
        "export_format_id": format,  
    }  
    response = clients['model'].export_model(name=name, output_config=output_config)  
    print("Long running operation:", response.operation.name)  
    result = response.result(timeout=1800)  
    metadata = response.operation.metadata  
    artifact_uri = str(metadata.value).split("\\\\")[1][4:-1]  
    print("Artifact Uri", artifact_uri)  
    return artifact_uri  
  
model_package = export_model(model_to_deploy_id, "tflite", MODEL_DIR)
```

Specify format and GCS location to export the edge packaged model artifacts.

## Image Model Exported to Edge - TFLite Interpreter


```
import tensorflow as tf

interpreter = tf.lite.Interpreter(model_path=tfllite_path)
interpreter.allocate_tensors()

input_details = interpreter.get_input_details()
output_details = interpreter.get_output_details()
input_shape = input_details[0]['shape']

print("input tensor shape", input_shape)
```

Instantiate TFLite interpreter for edge model.



TFLite run-time environment must be installed on edge device.

Run-time is smaller than TF run-time to fit into smaller memory.


## Image Model Exported to Edge - Image Resizing

```
test_items = ! gsutil cat $IMPORT_FILE | head -n1
test_item = test_items[0].split(',')[0]

with tf.io.gfile.GFile(test_item, "rb") as f:
    content = f.read()
test_image = tf.io.decode_jpeg(content)
print("test image shape", test_image.shape)

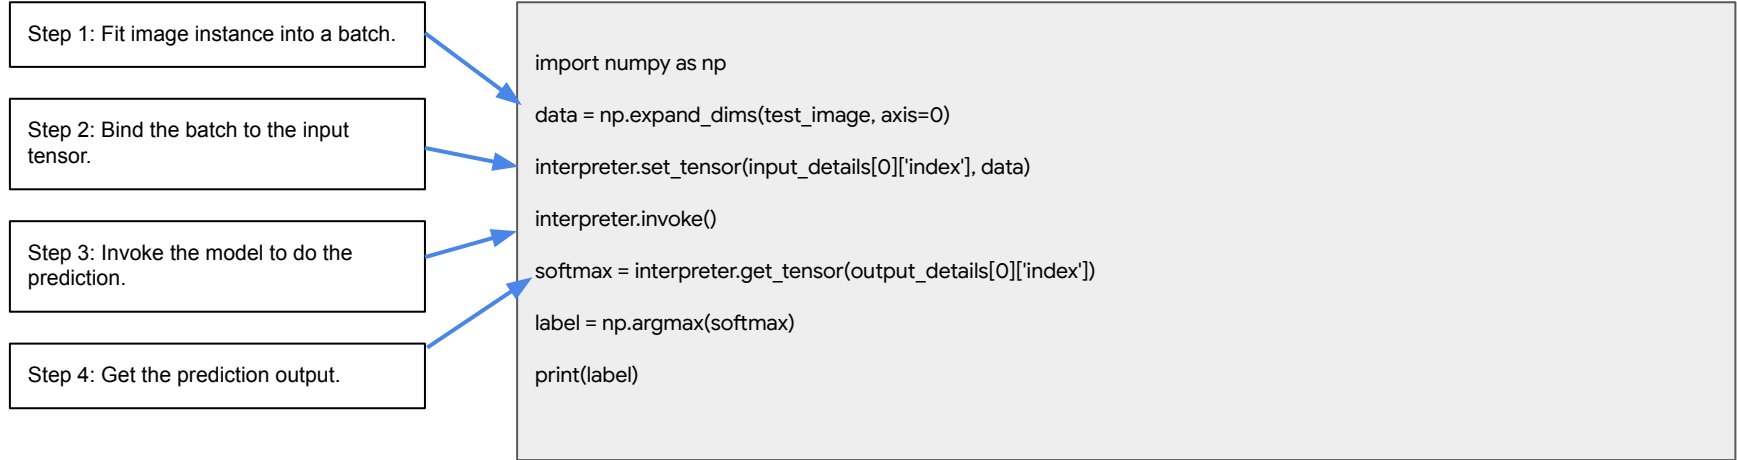
test_image = tf.image.resize(test_image, (224, 224))
print("test image shape", test_image.shape, test_image.dtype)

test_image = tf.cast(test_image, dtype=tf.uint8).numpy()
```



Must resize the image to the edge model input size, either upstream or on edge device.

## Image Model Exported to Edge - Prediction



# Workshop 3: Text Models


- Text Classification
- Text Sentiment Analysis
- Text Entity Extraction

# Workshop 3: AutoML Text Classification

## Text Classification (TCN) - Schema

```
# Text Dataset type  
DATA_SCHEMA = 'google-cloud-aiplatform/schema/dataset/metadata/text_1.0.0.yaml'  
# Text Labeling type  
LABEL_SCHEMA =  
"gs://google-cloud-aiplatform/schema/dataset/ioformat/text_classification_single_label_io_format_1.0.0.yaml"  
# Text Training task  
TRAINING_SCHEMA = "gs://google-cloud-aiplatform/schema/trainingjob/definition/automl_text_classification_1.0.0.yaml"
```

DATA specific to Text  
LABEL and TRAINING SCHEMA  
specific to TCN

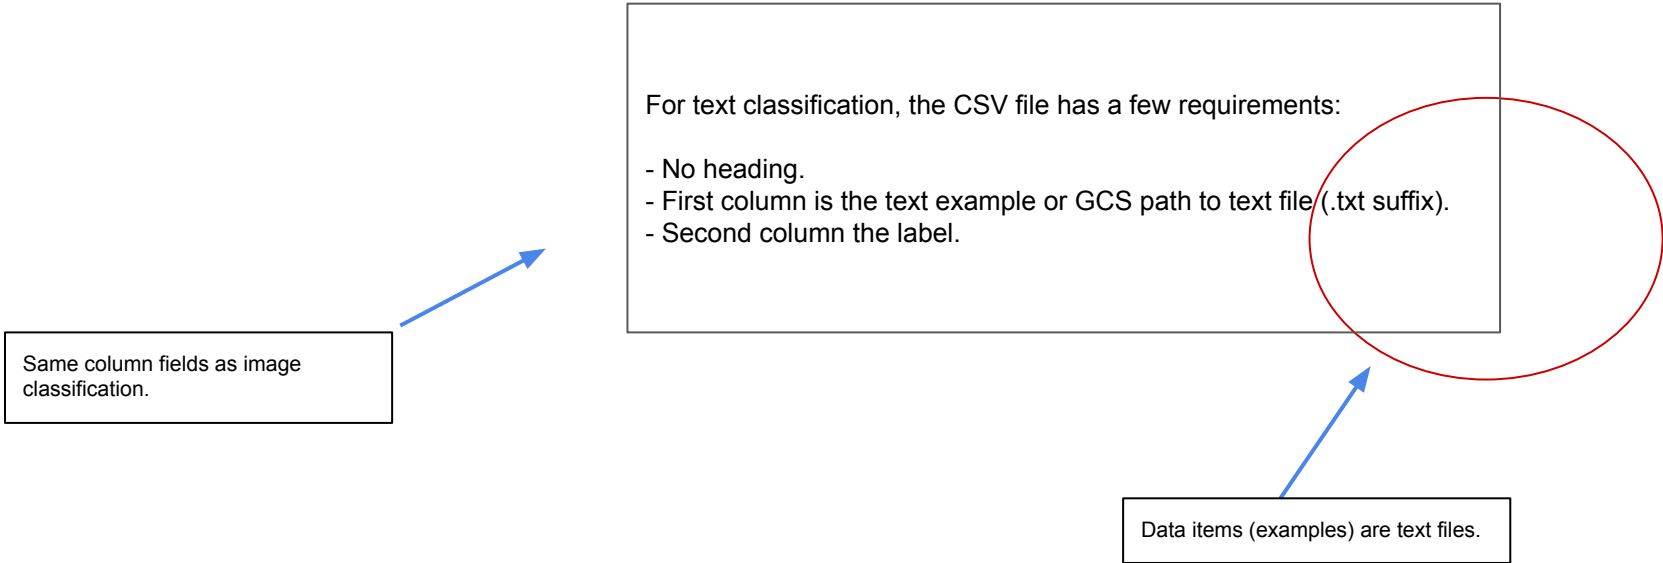


## Text Classification - Labeling

For text classification, the CSV file has a few requirements:

- No heading.
- First column is the text example or GCS path to text file (.txt suffix).
- Second column the label.

Same column fields as image classification.



```
graph LR; A[Same column fields as image classification.] --> B[For text classification, the CSV file has a few requirements:]; C[Data items (examples) are text files.] --> B; B --- D(( )); D --- E[First column is the text example or GCS path to text file (.txt suffix).]; D --- F[Second column the label.];
```

Data items (examples) are text files.



## Text Classification (TCN) - Task Requirements

Cloud only model.  
Can pick between single or  
multi-label classification.

```
PIPE_NAME = "happydb_pipe-" + TIMESTAMP  
MODEL_NAME = "happydb_model-" + TIMESTAMP
```

```
task = json_format.ParseDict({'multi_label': False,  
                             }, Value())
```

```
response = create_pipeline(PIPE_NAME, MODEL_NAME, dataset_id, TRAINING_SCHEMA, task)
```

## Text Classification (TCN) - Prediction

Format:

`{ 'content': text_item }`

Either text example, or GCS path to text file.

The ``response`` object returns a list, where each element in the list corresponds to the corresponding text item in the request. You will see in the output for each prediction:

- - Confidence level in the prediction (``confidences``).
- - The predicted label (``displayNames``).

Same as image classification

## Text Classification - Batch Prediction


For JSONL file, you make one dictionary entry per line for each data item (instance). The dictionary contains the key/value pairs:

- `content`: The Cloud Storage path to the file with the text item.
- `mime\_type`: The content type. In our example, it is an `text` file.

For example:

`{'content': '[your-bucket]/file1.txt', 'mime_type': 'text'}`

Same as image model, except  
content is text file.



# Workshop 3: AutoML Text Sentiment Analysis

## Text Sentiment Analysis (TST) - Schema

```
# Text Dataset type
DATA_SCHEMA = 'google-cloud-aiplatform/schema/dataset/metadata/text_1.0.0.yaml'
# Text Labeling type
LABEL_SCHEMA = "gs://google-cloud-aiplatform/schema/dataset/ioformat/text_sentiment_io_format_1.0.0.yaml"
# Text Training task
TRAINING_SCHEMA = "gs://google-cloud-aiplatform/schema/trainingjob/definition/automl_text_sentiment_1.0.0.yaml"
```

LABEL and TRAINING SCHEMA  
specific to TST

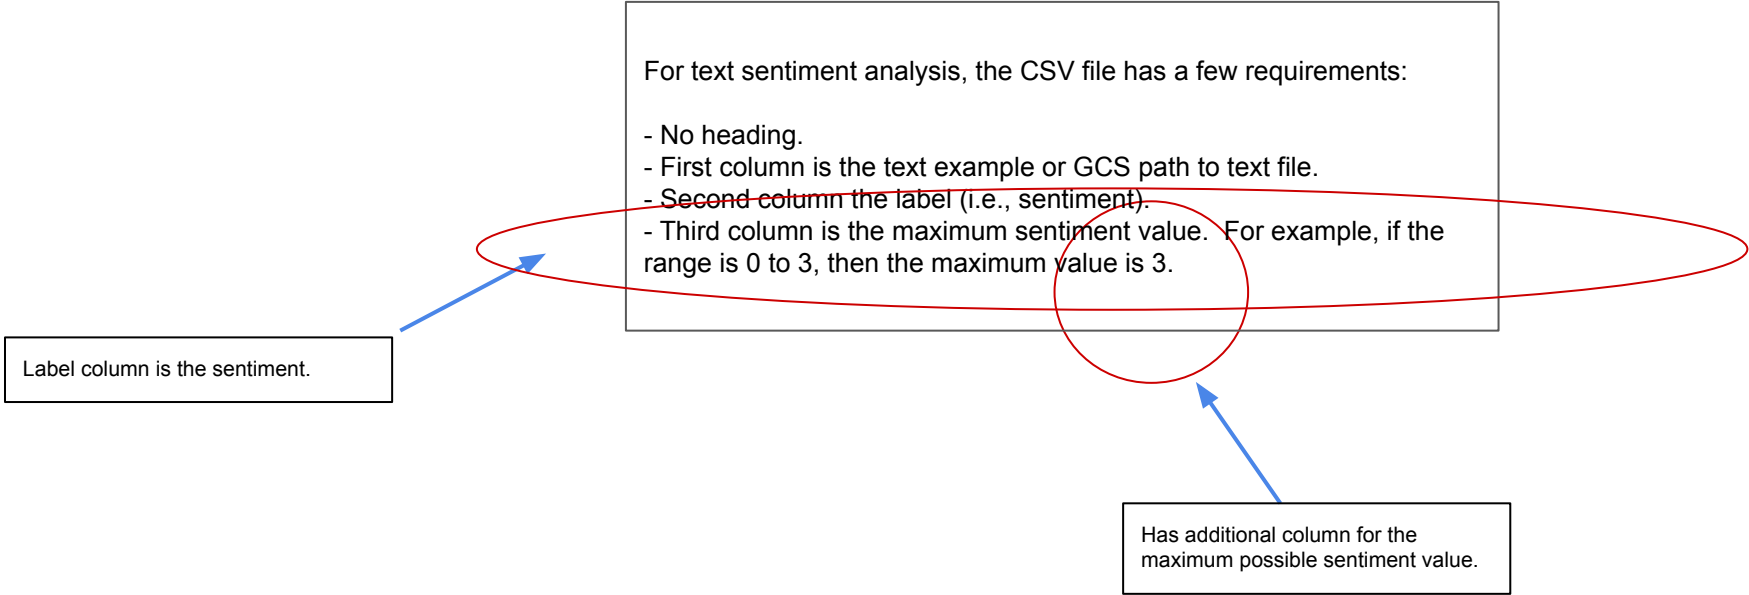


## Text Sentiment Analysis - Labeling

For text sentiment analysis, the CSV file has a few requirements:

- No heading.
- First column is the text example or GCS path to text file.
- ~~Second column the label (i.e., sentiment).~~
- Third column is the maximum sentiment value. For example, if the range is 0 to 3, then the maximum value is 3.

Label column is the sentiment.



Has additional column for the maximum possible sentiment value.

## Text Sentiment Analysis (TST) - Task Requirements

```
PIPE_NAME = "claritin_pipe-" + TIMESTAMP  
MODEL_NAME = "claritin_model-" + TIMESTAMP
```

```
task = json_format.ParseDict({'sentiment_max': SENTIMENT_MAX,  
                             }, Value())
```

```
response = create_pipeline(PIPE_NAME, MODEL_NAME, dataset_id, TRAINING_SCHEMA, task)
```

Cloud only model.

Specify the maximum sentiment.



## Text Sentiment Analysis - Prediction

Format:

{ 'content': text\_item }

Same as text classification

The response object returns a list, where each element in the list corresponds to the corresponding image in the request. You will see in the output for each prediction -- in our case there is just one:

- The sentiment rating

The sentiment rating



## Text Sentiment Analysis - Batch Prediction


For JSONL file, you make one dictionary entry per line for each data item (instance). The dictionary contains the key/value pairs:

- `content`: The Cloud Storage path to the file with the text item.
- `mime\_type`: The content type. In our example, it is an `text` file.

For example:

`{'content': '[your-bucket]/file1.txt', 'mime_type': 'text'}`

Same as text classification



# Workshop 3: AutoML Text Entity Extraction

## Text Entity Extraction (TEN) - Schema

```
# Text Dataset type
DATA_SCHEMA = 'google-cloud-aiplatform/schema/dataset/metadata/text_1.0.0.yaml'
# Text Labeling type
LABEL_SCHEMA = "gs://google-cloud-aiplatform/schema/dataset/ioformat/text_extraction_io_format_1.0.0.yaml"
# Text Training task
TRAINING_SCHEMA = "gs://google-cloud-aiplatform/schema/trainingjob/definition/automl_text_extraction_1.0.0.yaml"
```

LABEL and TRAINING SCHEMA  
specific to TEN

## Text Entity Extraction - Labeling

For text entity extraction, the JSONL file has a few requirements:

- Each data item is a separate JSON object, on a separate line.
- The key/value pair `text\_segment\_annotations` is a list of character start/end positions in the text per entity with the corresponding label.
- `display\_name`: ~~The label.~~
- `start\_offset/end\_offset`: The character offsets of the start/end of the entity.
- The key/value pair `text\_content` is the text.

For example:

```
{'text_segment_annotations': [{'end_offset': value, 'start_offset': value, 'display_name': label}, ...], 'text_content': text}
```

Each entity is specified with a start and end position in the text.

## Text Entity Extraction (TEN) - Task Requirements

```
task = json_format.ParseDict({'multi_label': False,  
    'budget_milli_node_hours': 8000,  
    'model_type': "CLOUD",  
    'disable_early_stopping': False  
}, Value())
```

Cloud only.  
Entities can have multiple labels.

## Text Entity Extraction - Prediction

Format:

{ 'content': text\_item }

Same as text classification

The `response` object returns a list, where each element in the list corresponds to the corresponding data item in the request. You will see in the output for each prediction -- in our case there is just one:

- `prediction`: A list of IDs assigned to each entity extracted from the text.
- `confidences`: The confidence level between 0 and 1 for each entity.
- `display\_names`: The label name for each entity.
- `textSegmentStartOffsets`: The character start location of the entity in the text.
- `textSegmentEndOffsets`: The character end location of the entity in the text.

The location of each entity, label of each entity, and confidence score

## Text Entity Extraction - Batch Prediction

For JSONL file, you make one dictionary entry per line for each data item (instance). The dictionary contains the key/value pairs:

- `content`: The Cloud Storage path to the file with the text item.
- `mime\_type`: The content type. In our example, it is an `text` file.

For example:

`{'content': '[your-bucket]/file1.txt', 'mime_type': 'text'}`

Same as text classification and sentiment analysis.

